



Applied Math for Deep Learning

Prof. Kuan-Ting Lai


2020/3/10

Applied Math for Deep Learning

- Linear Algebra
- Probability
- Calculus
- Optimization

Linear Algebra



- Scalar
 - real numbers
- Vector (1D)
 - Has a magnitude & a direction
- Matrix (2D)
 - An array of numbers arranged in rows & columns
- Tensor ($\geq 3D$)
 - Multi-dimensional arrays of numbers



$$\mathbf{A} = \begin{pmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{pmatrix}$$

Real-world examples of Data Tensors

- Timeseries Data – 3D (samples, timesteps, features)
- Images – 4D (samples, height, width, channels)
- Video – 5D (samples, frames, height, width, channels)



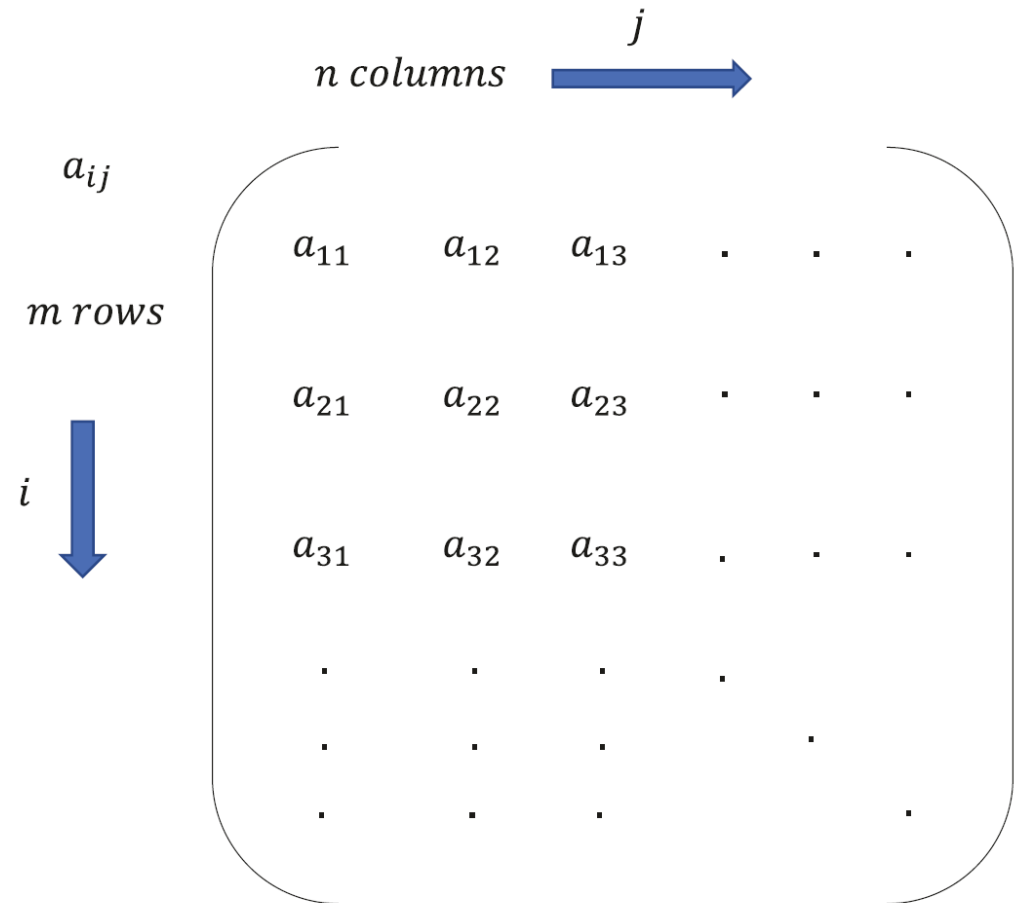
A full-page image of Keanu Reeves as Neo from the movie The Matrix. He is wearing a black turtleneck and dark sunglasses, looking directly at the camera. The background is a green digital rain effect, with white text resembling code or data floating around him. The title 'The Matrix' is written in large white letters across the bottom center of the image.

The Matrix

Matrix

- Define a matrix with m rows and n columns:

$$A_{m \times n} \in \mathbb{R}^{m \times n}$$



Matrix Operations

- Addition and Subtraction

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \quad B = \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix}$$

$$A + B = \begin{bmatrix} 1+5 & 2+6 \\ 3+7 & 4+8 \end{bmatrix} = \begin{bmatrix} 6 & 8 \\ 10 & 12 \end{bmatrix}$$


$$A - B = \begin{bmatrix} 1-5 & 2-6 \\ 3-7 & 4-8 \end{bmatrix} = \begin{bmatrix} -4 & -4 \\ -4 & -4 \end{bmatrix}$$

Matrix Multiplication

- Two matrices A and B, where $A \in \mathbb{R}^{m \times n}$ $B \in \mathbb{R}^{p \times q}$
- The columns of A must be equal to the rows of B, i.e. $n == p$

- $A * B = C$, where $C \in \mathbb{R}^{m \times q}$

- $$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$



Example of Matrix Multiplication (3-1)

"Dot Product"

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 & \\ & \end{bmatrix}$$

Example of Matrix Multiplication (3-2)

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 & 64 \\ & \end{bmatrix}$$

Example of Matrix Multiplication (3-3)

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 & 64 \\ 139 & 154 \end{bmatrix} \checkmark$$

Matrix Transpose


$$A \in \mathbb{R}^{m \times n} \quad A^T \in \mathbb{R}^{n \times m}$$

$$a'_{ji} = a_{ij} \quad \forall i \in \{1, 2, \dots, m\}, \forall j \in \{1, 2, \dots, n\}$$

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{bmatrix}$$

A

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{bmatrix}$$



A^T

$$\begin{bmatrix} 1 & 3 & 5 \\ 2 & 4 & 6 \end{bmatrix}$$

Dot Product


- Dot product of two vectors become a **scalar**
- Notation: $v_1 \cdot v_2$ or $v_1^T v_2$

$$v_1 = \begin{bmatrix} v_{11} \\ v_{12} \\ \cdot \\ \cdot \\ \cdot \\ v_{1n} \end{bmatrix} \quad v_2 = \begin{bmatrix} v_{21} \\ v_{22} \\ \cdot \\ \cdot \\ \cdot \\ v_{2n} \end{bmatrix}$$

$$v_1 \cdot v_2 = v_1^T v_2 = v_2^T v_1 = v_{11}v_{21} + v_{12}v_{22} + \dots + v_{1n}v_{2n} = \sum_{k=1}^n v_{1k}v_{2k}$$

Dot Product of Matrix

$$x \cdot y = z$$




Linear Independence

- A vector is **linearly dependent** on other vectors if it can be expressed as the linear combination of other vectors
- A set of vectors v_1, v_2, \dots, v_n is **linearly independent** if $a_1 v_1 + a_2 v_2 + \dots + a_n v_n = 0$ implies all $a_i = 0, \forall i \in \{1, 2, \dots, n\}$

$$\begin{bmatrix} v_1 & v_2 & \dots & v_n \end{bmatrix} \begin{bmatrix} a_1 \\ a_2 \\ \cdot \\ \cdot \\ a_n \end{bmatrix} = 0 \text{ where } v_i \in \mathbb{R}^{m \times 1} \forall i \in \{1, 2, \dots, n\}, \begin{bmatrix} a_1 \\ a_2 \\ \cdot \\ \cdot \\ a_n \end{bmatrix} \in \mathbb{R}^{n \times 1}$$

Span the Vector Space

- n linearly independent vectors can span n -dimensional space



Rank of a Matrix

- Rank is:
 - The number of linearly independent row or column vectors
 - The dimension of the vector space generated by its columns
- Row rank = Column rank
- Example:

$$\mathbf{A} = \begin{bmatrix} \mathbf{1} & \mathbf{2} & \mathbf{1} \\ \mathbf{-2} & \mathbf{-3} & \mathbf{1} \\ \mathbf{3} & \mathbf{5} & \mathbf{0} \end{bmatrix} \xrightarrow{\text{Row-echelon form}} \begin{bmatrix} 1 & 0 & -5 \\ 0 & 1 & 3 \\ 0 & 0 & 0 \end{bmatrix}$$

Identity Matrix I

- Any vector or matrix multiplied by I remains unchanged
- For a matrix $A_{m \times n}$, $AI_n = I_mA = A$

$$I = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \in \mathbb{R}^{3 \times 3}$$

$$Iv = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 2 \\ 3 \\ 4 \end{bmatrix} = \begin{bmatrix} 2 \\ 3 \\ 4 \end{bmatrix}$$

Inverse of a Matrix

- The product of a **square** matrix A and its inverse matrix A^{-1} produces the identity matrix I
- $AA^{-1} = A^{-1}A = I$
- Inverse matrix is square, but not all square matrices has inverses

Pseudo Inverse

- Non-square matrix and have left-inverse or right-inverse matrix
- Example:

$$Ax = b, A \in \mathbb{R}^{m \times n}, b \in \mathbb{R}^n$$

- Create a square matrix $A^T A$

$$A^T Ax = A^T b$$

- Multiplied both sides by inverse matrix $(A^T A)^{-1}$

$$x = (A^T A)^{-1} A^T b$$

- $(A^T A)^{-1} A^T$ is the pseudo inverse function

Norm

- Norm is a measure of a vector's magnitude

- l_2 norm
$$\|x\|_2 = \left(|x_1|^2 + |x_2|^2 + \dots + |x_n|^2 \right)^{1/2} = (x \cdot x)^{1/2} = (x^T x)^{1/2}$$

- l_1 norm
$$\|x\|_1 = |x_1| + |x_2| + \dots + |x_n|$$




- l_p norm
$$\left(|x_1|^p + |x_2|^p + \dots + |x_n|^p \right)^{1/p}$$

- l_∞ norm

$$\lim_{p \rightarrow \infty} \|x\|_p = \lim_{p \rightarrow \infty} \left(|x_1|^p + |x_2|^p + \dots + |x_n|^p \right)^{1/p} = \max(x_1, x_2, \dots, x_n)$$

Unit norms in 2D Vectors

- The set of all vectors of norm 1 in different 2D norms




L1 and L2 Regularization

$$\text{Min } \sum_{i=1}^N (y_i - \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i})^2$$

subject to

$$\sum_{j=1}^2 \beta_j^2 \leq s$$



LASSO

RIDGE REGRESSION

Eigen Vectors



- Eigenvector is a non-zero vector that changed by only a scalar factor λ when linear transform A is applied to:

$$Ax = \lambda x, A \in \mathbb{R}^{n \times n}, x \in \mathbb{R}^n$$

- x are Eigenvectors and λ are Eigenvalues
- One of the most important concepts for machine learning, ex:
 - Principle Component Analysis (PCA)
 - Eigenvector centrality
 - PageRank
 - ...

Example: Shear Mapping


- Horizontal axis is the Eigenvector



Power Iteration Method for Computing Eigenvector

1. Start with random vector v
2. Calculate iteratively: $v^{(k+1)} = A^k v$
3. After v^k converges, $v^{(k+1)} \cong v^k$
4. v^k will be the Eigenvector with largest Eigenvalue

NumPy for Linear Algebra



- NumPy is the fundamental package for scientific computing with Python. It contains among other things:
 - a powerful N-dimensional array object
 - sophisticated (broadcasting) functions
 - tools for integrating C/C++ and Fortran code
 - useful linear algebra, Fourier transform, and random number capabilities

Python & NumPy tutorial

- <http://cs231n.github.io/python-numpy-tutorial/>
- Stanford CS231n: Convolutional Neural Networks for Visual Recognition
 - <http://cs231n.stanford.edu/>

Table of contents:

- Python
 - Basic data types
 - Containers
 - Lists
 - Dictionaries
 - Sets
 - Tuples
 - Functions
 - Classes
- Numpy
 - Arrays
 - Array indexing
 - Datatypes
 - Array math
 - Broadcasting
- SciPy
 - Image operations
 - MATLAB files
 - Distance between points
- Matplotlib
 - Plotting
 - Subplots
 - Images

Create Tensors

Scalars (0D tensors)

```
>>> import numpy as np
>>> x = np.array(12)
>>> x
array(12)
>>> x.ndim
0
```

Vectors (1D tensors)

```
>>> x = np.array([12, 3, 6, 14])
>>> x
array([12, 3, 6, 14])
>>> x.ndim
1
```

Matrices (2D tensors)

```
>>> x = np.array([[5, 78, 2, 34, 0],
                  [6, 79, 3, 35, 1],
                  [7, 80, 4, 36, 2]])
>>> x.ndim
2
```

Create 3D Tensor

```
>>> x = np.array([[[5, 78, 2, 34, 0],  
                  [6, 79, 3, 35, 1],  
                  [7, 80, 4, 36, 2]],  
                 [[5, 78, 2, 34, 0],  
                  [6, 79, 3, 35, 1],  
                  [7, 80, 4, 36, 2]],  
                 [[5, 78, 2, 34, 0],  
                  [6, 79, 3, 35, 1],  
                  [7, 80, 4, 36, 2]])  
  
>>> x.ndim  
3
```


Attributes of a Tensor

- Number of axes (dimensions)
 - `x.ndim`
- Shape
 - This is a tuple of integers showing how many data the tensor has along each axis
- Data type
 - `uint8`, `float32` or `float64`

Manipulating Tensors in Numpy

```
>>> my_slice = train_images[10:100]
>>> print(my_slice.shape)
(90, 28, 28)
```

```
>>> my_slice = train_images[10:100, :, :] ← Equivalent to the
>>> my_slice.shape                               previous example
(90, 28, 28)
>>> my_slice = train_images[10:100, 0:28, 0:28] ← Also equivalent to the
>>> my_slice.shape                               previous example
(90, 28, 28)
```

```
my_slice = train_images[:, 7:-7, 7:-7]
```

Displaying the Fourth Digit

```
digit = train_images[4]  
  
import matplotlib.pyplot as plt  
plt.imshow(digit, cmap=plt.cm.binary)  
plt.show()
```






Figure 2.2 The fourth sample in our dataset

Real-world examples of Data Tensors

- Vector data – 2D (samples, features)
- Timeseries Data – 3D (samples, timesteps, features)
- Images – 4D (samples, height, width, channels)
- Video – 5D (samples, frames, height, width, channels)



Batch size & Epochs

- A sample
 - A sample is a single row of data
- Batch size
 - Number of samples used for one iteration of gradient descent
 - Batch size = 1: stochastic gradient descent
 - $1 < \text{Batch size} < \text{all}$: mini-batch gradient descent
 - Batch size = all: batch gradient descent
- Epoch
 - Number of times that the learning algorithm work through all training samples

Element-wise Operations for Matrix

- Operate on each element

```
def naive_add(x, y):  
    assert len(x.shape) == 2  
    assert x.shape == y.shape
```

← **x and y are 2D
Numpy tensors.**

```
    x = x.copy()  
    for i in range(x.shape[0]):  
        for j in range(x.shape[1]):  
            x[i, j] += y[i, j]  
    return x
```

← **Avoid overwriting
the input tensor.**

NumPy Operation for Matrix

- Leverage the Basic Linear Algebra subprograms (BLAS)
- BLAS is optimized using C or Fortran

```
import numpy as np
```

```
z = x + y
```

← **Element-wise addition**

```
z = np.maximum(z, 0.)
```

← **Element-wise relu**

Broadcasting

- Apply smaller tensor repeated to the extra axes of the larger tensor

```
def naive_add_matrix_and_vector(x, y):  
    assert len(x.shape) == 2  
    assert len(y.shape) == 1  
    assert x.shape[1] == y.shape[0]  
  
    x = x.copy()  
    for i in range(x.shape[0]):  
        for j in range(x.shape[1]):  
            x[i, j] += y[j]  
    return x
```

x is a 2D Numpy tensor.


y is a Numpy vector.

**Avoid overwriting
the input tensor.**

Tensor Dot

```
import numpy as np
```

```
z = np.dot(x, y)
```



Implementation of Dot Product

x and y are Numpy matrices.

```
def naive_matrix_dot(x, y):  
    assert len(x.shape) == 2  
    assert len(y.shape) == 2  
    assert x.shape[1] == y.shape[0]  
  
    z = np.zeros((x.shape[0], y.shape[1]))  
    for i in range(x.shape[0]):  
        for j in range(y.shape[1]):  
            row_x = x[i, :]  
            column_y = y[:, j]  
            z[i, j] = naive_vector_dot(row_x, column_y)  
    return z
```

← The first dimension of x must be the same as the 0th dimension of y!

← This operation returns a matrix of 0s with a specific shape.

← Iterates over the rows of x ...

← ... and over the columns of y.

Tensor Reshaping

- Rearrange a tensor's rows and columns to match a target shape

```
>>> x = np.array([[0., 1.],  
                  [2., 3.],  
                  [4., 5.]])
```

```
>>> print(x.shape)  
(3, 2)
```

```
>>> x = x.reshape((6, 1))
```

```
>>> x  
array([[ 0.],  
       [ 1.],  
       [ 2.],  
       [ 3.],  
       [ 4.],  
       [ 5.]])
```

```
>>> x = x.reshape((2, 3))
```

```
>>> x  
array([[ 0.,  1.,  2.],  
       [ 3.,  4.,  5.]])
```

Matrix Transposition

- Transposing a matrix means exchanging its rows and its columns

```
>>> x = np.zeros((300, 20))  
>>> x = np.transpose(x)  
>>> print(x.shape)  
(20, 300)
```

← **Creates an all-zeros matrix
of shape (300, 20)**

Unfolding the Manifold

- Tensor operations are complex geometric transformation in high-dimensional space
 - Dimension reduction



Maria Ghetana Agnesi

$$(\ln x)' = \frac{1}{x} \quad \int \frac{1}{x} dx = \ln|x| + C$$

$$f(x) = x^2$$
$$\int \sin x dx = -\cos x + C$$

$$\int_a^b f'(x) dx = f(b) - f(a)$$

$$m \frac{d^2 x}{dt^2} = -kx$$

$$\frac{df(x)}{dz}$$

Calculus

$$x^2 - 3x - 4 = 0$$


$$4x^2 - 3x - 1 = 0$$

$$\int f(x) dx$$



$$\frac{dA}{dt} = \frac{dB}{dt} = -\frac{dC}{dt} = -\frac{dD}{dt} = (d_1)T^{\frac{1}{2}}AB - (d_2)T^{\frac{1}{2}}CD$$

$$x^2 = A \quad \frac{dT}{dt} = (c_3) \frac{dA}{dt} - (c_4)(T_0 - T)$$




Differentiation

$$\frac{df}{dt} = \lim_{h \rightarrow 0} \frac{f(t+h) - f(t)}{h}$$

OR

$$\frac{df}{dt} = \lim_{h \rightarrow 0} \frac{f(t+h) - f(t-h)}{2h}$$



Gradient of a Function

- Gradient is a multi-variable generalization of the derivative
- Apply partial derivatives

$$\nabla f = \left[\frac{\partial f}{\partial x_1} \frac{\partial f}{\partial x_2} \dots \frac{\partial f}{\partial x_n} \right]^T$$

- Example

$$f(x, y, z) = x + y^2 + z^3$$

$$\nabla f = \left[1 \ 2y \ 3z^2 \right]^T$$

Hessian Matrix

- Second-order partial derivatives

$$Hf = \begin{bmatrix} \frac{\delta^2 f}{\delta x^2} & \frac{\delta^2 f}{\delta x \delta y} & \frac{\delta^2 f}{\delta x \delta z} \\ \frac{\delta^2 f}{\delta y \delta x} & \frac{\delta^2 f}{\delta y^2} & \frac{\delta^2 f}{\delta y \delta z} \\ \frac{\delta^2 f}{\delta z \delta x} & \frac{\delta^2 f}{\delta z \delta y} & \frac{\delta^2 f}{\delta z^2} \end{bmatrix}$$




Maxima and Minima for Univariate Function

- If $\frac{df(x)}{dx} = 0$, it's a minima or a maxima point, then we study the second derivative:

– If $\frac{d^2f(x)}{dx^2} < 0 \Rightarrow$ Maxima

– If $\frac{d^2f(x)}{dx^2} > 0 \Rightarrow$ Minima

– If $\frac{d^2f(x)}{dx^2} = 0 \Rightarrow$ Point of reflection



Maxima and Minima for Multivariate Function

- Computing the gradient and setting it to zero vector would give us the list of stationary points.
- For a stationary point $x_0 \in \mathbb{R}^n$
 - If the Hessian matrix of the function at x_0 has both positive and negative eigen values, then x_0 is a saddle point
 - If the eigen values of the Hessian matrix are all positive then the stationary point is a local minima
 - If the eigen values are all negative then the stationary point is a local maxima

Chain Rule

$$\frac{dy}{dx} = \frac{dy}{du} \frac{du}{dx}$$

$$\frac{d^2 y}{dx^2} = \frac{d^2 y}{du^2} \left(\frac{du}{dx} \right)^2 + \frac{dy}{du} \frac{d^2 u}{dx^2}$$

$$\frac{d^3 y}{dx^3} = \frac{d^3 y}{du^3} \left(\frac{du}{dx} \right)^3 + 3 \frac{d^2 y}{du^2} \frac{du}{dx} \frac{d^2 u}{dx^2} + \frac{dy}{du} \frac{d^3 u}{dx^3}$$

$$\frac{d^4 y}{dx^4} = \frac{d^4 y}{du^4} \left(\frac{du}{dx} \right)^4 + 6 \frac{d^3 y}{du^3} \left(\frac{du}{dx} \right)^2 \frac{d^2 u}{dx^2} + \frac{d^2 y}{du^2} \left(4 \frac{du}{dx} \frac{d^3 u}{dx^3} + 3 \left(\frac{d^2 u}{dx^2} \right)^2 \right) + \frac{dy}{du} \frac{d^4 u}{dx^4}.$$

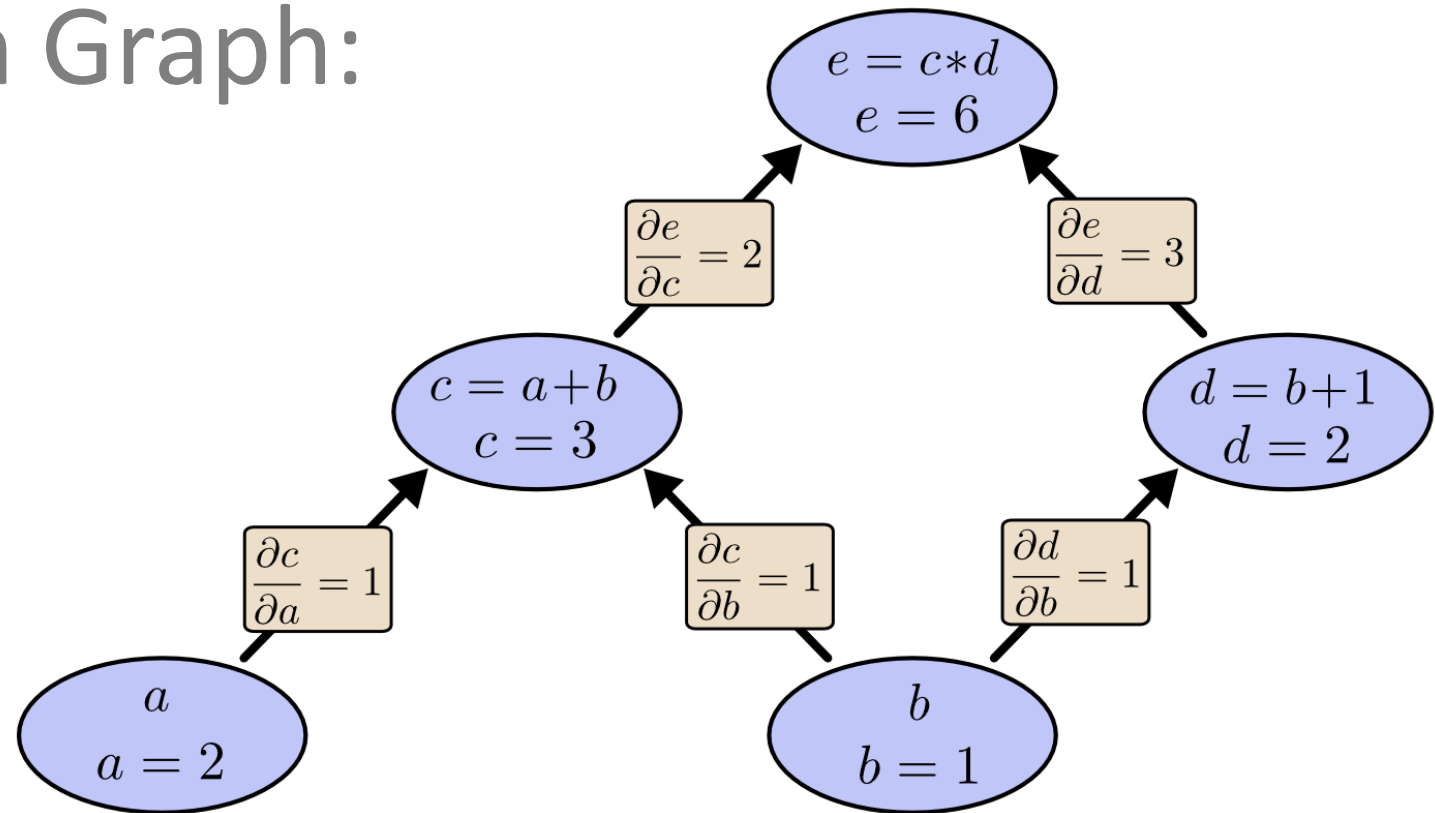
Symbolic Differentiation

Computation Graph:


$$c = a + b$$

$$d = b + 1$$

$$e = c * d$$




Stochastic Gradient Descent



1. Draw a batch of training samples x and corresponding targets y
2. Run the network on x to obtain predictions y_{pred}
3. Compute the loss of the network on the batch, a measure of the mismatch between y_{pred} and y
4. Compute the gradient of the loss with regard to the network's parameters (a backward pass).
5. Move the parameters a little in the opposite direction from the gradient: $W -= \text{step} * \text{gradient}$

Gradient Descent along a 2D Surface




Avoid Local Minimum using Momentum

```
past_velocity = 0.  
momentum = 0.1  
while loss > 0.01:  
    w, loss, gradient = get_current_parameters()  
    velocity = past_velocity * momentum + learning_rate * gradient  
    w = w + momentum * velocity - learning_rate * gradient  
    past_velocity = velocity  
    update_parameter(w)
```

Constant momentum factor

Optimization loop





Basics of Probability

Three Axioms of Probability


- Given an Event E in a sample space S , $S = \bigcup_{i=1}^N E_i$
- First axiom
 - $P(E) \in \mathbb{R}, 0 \leq P(E) \leq 1$
- Second axiom
 - $P(S) = 1$
- Third axiom
 - Additivity, any countable sequence of mutually exclusive events E_i
 - $P(\bigcup_{i=1}^n E_i) = P(E_1) + P(E_2) + \cdots + P(E_n) = \sum_{i=1}^n P(E_i)$

Union, Intersection, and Conditional Probability

- $P(A \cup B) = P(A) + P(B) - P(A \cap B)$
- $P(A \cap B)$ is simplified as $P(AB)$
- Conditional Probability $P(A|B)$, the probability of event A given B has occurred

$$- P(A|B) = P\left(\frac{AB}{B}\right)$$

$$- P(AB) = P(A|B)P(B) = P(B|A)P(A)$$



Chain Rule of Probability

- The joint probability can be expressed as chain rule

$$\begin{aligned} P(A_1 A_2 A_3 \dots A_n) &= P(A_1) P(A_2 / A_1) P(A_3 / A_1 A_2) \dots P(A_n / A_1 A_2 \dots A_{(n-1)}) \\ &= P(A_1) \prod_{i=2}^n P(A_i / A_1 A_2 A_3 \dots A_{(n-1)}) \end{aligned}$$

Mutually Exclusive

- $P(AB) = 0$
- $P(A \cup B) = P(A) + P(B)$

Independence of Events

- Two events A and B are said to be independent if the probability of their intersection is equal to the product of their individual probabilities

- $P(AB) = P(A)P(B)$

- $P(A|B) = P(A)$

Bayes Rule

- $P(A|B) = \frac{P(B|A)P(A)}{P(B)}$
- Proof:
 - Remember $P(A|B) = P\left(\frac{AB}{B}\right)$
 - So $P(AB) = P(A|B)P(B) = P(B|A)P(A)$
 - Then Bayes $P(A|B) = P(B|A)P(A)/P(B)$

Probability Mass Function and Dense Function

- Probability mass function (PMF)

- Function that gives the probability that a discrete random variable is exactly equal to some value

$$P(X = i) = \frac{1}{6}, i \in \{1, 2, 3, 4, 5, 6\}$$

- Probability dense function (PDF)

- Specify the probability of the random variable falling within a particular range of values

$$\int_D P(x) dx = 1$$

Expectation of a Random Variable

- Expectation of a discrete random variable

$$E[X] = x_1p_1 + x_2p_2 + \cdots + x_np_n = \sum_{i=1}^n x_ip_i$$

- Expectation of a continuous random variable

$$E[X] = \int_D xP(x)dx$$

Variance of a Random Variable

- Expectation of a discrete random variable

$$Var[X] = E[(X - \mu)^2], \text{ where } \mu = E[X]$$

- Expectation of a continuous random variable

$$Var[X] = \int_D (x - \mu)^2 P(x) dx$$

- Standard deviation σ is the square root of variance

Covariance and Correlation Coefficient

- Expectation of a discrete random variable

$$\text{Cov}(X, Y) = E[(X - \mu_x)(Y - \mu_y)],$$

where $\mu_x = E[X]$, $\mu_y = E[Y]$


- Correlation coefficient

$$\rho = \frac{\text{cov}(X, Y)}{\sigma_x \sigma_y}$$

Normal (Gaussian) Distribution

- One of the most important distributions
- Central limit theorem
 - Averages of samples of observations of random variables independently drawn from independent distributions converge to the normal distribution

$$f(x \mid \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



Optimization

The *standard form* of a continuous optimization problem is^[1]

$$\begin{array}{ll}\underset{x}{\text{minimize}} & f(x) \\ \text{subject to} & g_i(x) \leq 0, \quad i = 1, \dots, m \\ & h_j(x) = 0, \quad j = 1, \dots, p\end{array}$$

where

- $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is the **objective function** to be minimized over the n -variable vector x ,
- $g_i(x) \leq 0$ are called **inequality constraints**
- $h_j(x) = 0$ are called **equality constraints**, and
- $m \geq 0$ and $p \geq 0$.

Formulate Your Problem

- Linear model: $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$
- Least-squared Error: $(f(\mathbf{x}) - \mathbf{y})^2$
- Regularization: $\|\mathbf{w}\|$
- Objective function:

$$\min_{\mathbf{w}} (\mathbf{w}^T \mathbf{x} - \mathbf{y})^2 + \lambda \|\mathbf{w}\|$$

References

- Francois Chollet, “Deep Learning with Python,” Chapter 2 “Mathematical Building Blocks of Neural Networks”
- Santanu Pattanayak, “Pro Deep Learning with TensorFlow,” Apress, 2017
- [Machine Learning Cheat Sheet](#)
- <https://machinelearningmastery.com/difference-between-a-batch-and-an-epoch/>
- <https://www.quora.com/What-is-the-difference-between-L1-and-L2-regularization-How-does-it-solve-the-problem-of-overfitting-Which-regularizer-to-use-and-when>
- Wikipedia