Attention, Transformer and BERT

Prof. Kuan-Ting Lai

2020/6/16

Attention is All You Need!

A. Waswani et al., *NIPS*, 2017

Google Brain & University of Toronto

Attention

- Visual attention and textual attention

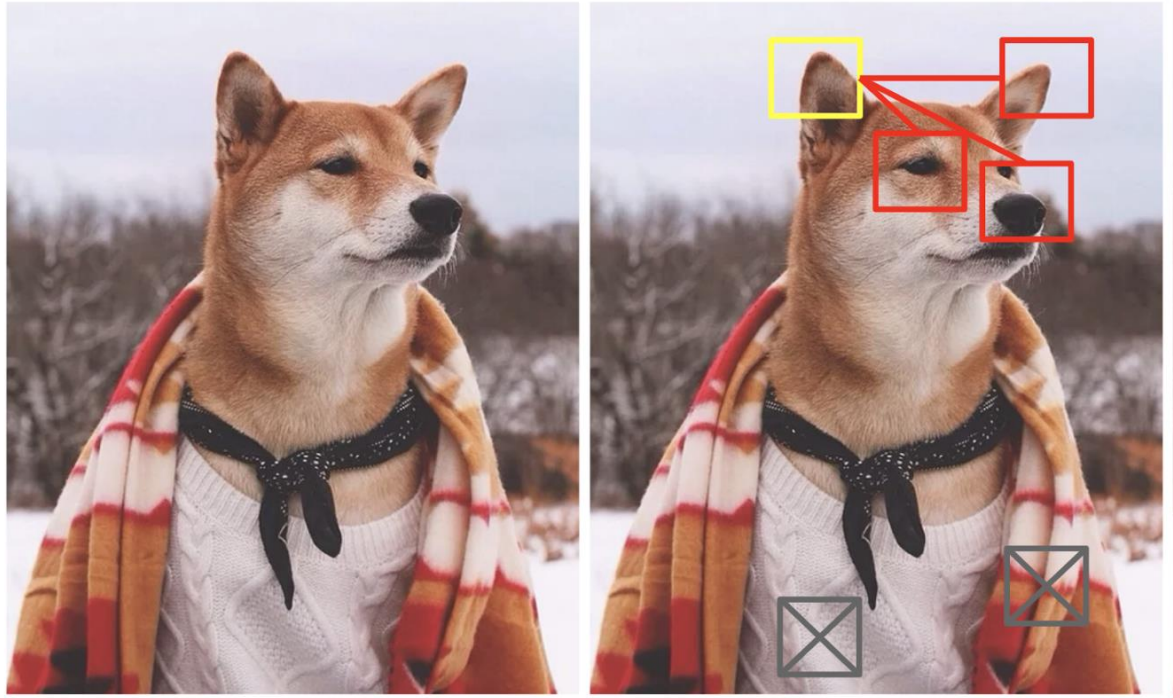

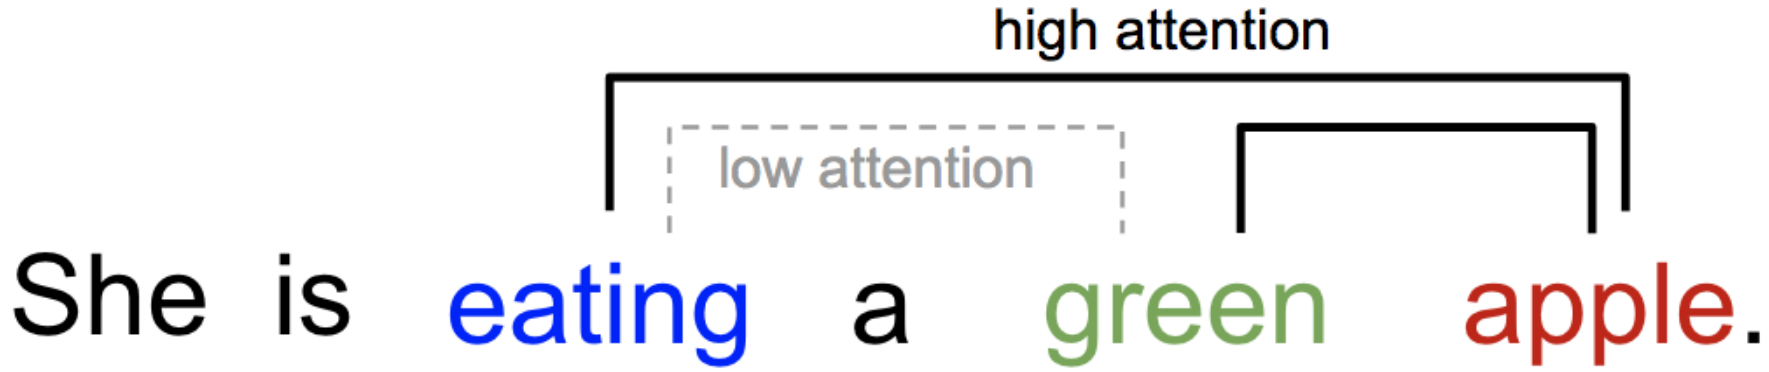

https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html

Seq2seq model

- Language translation

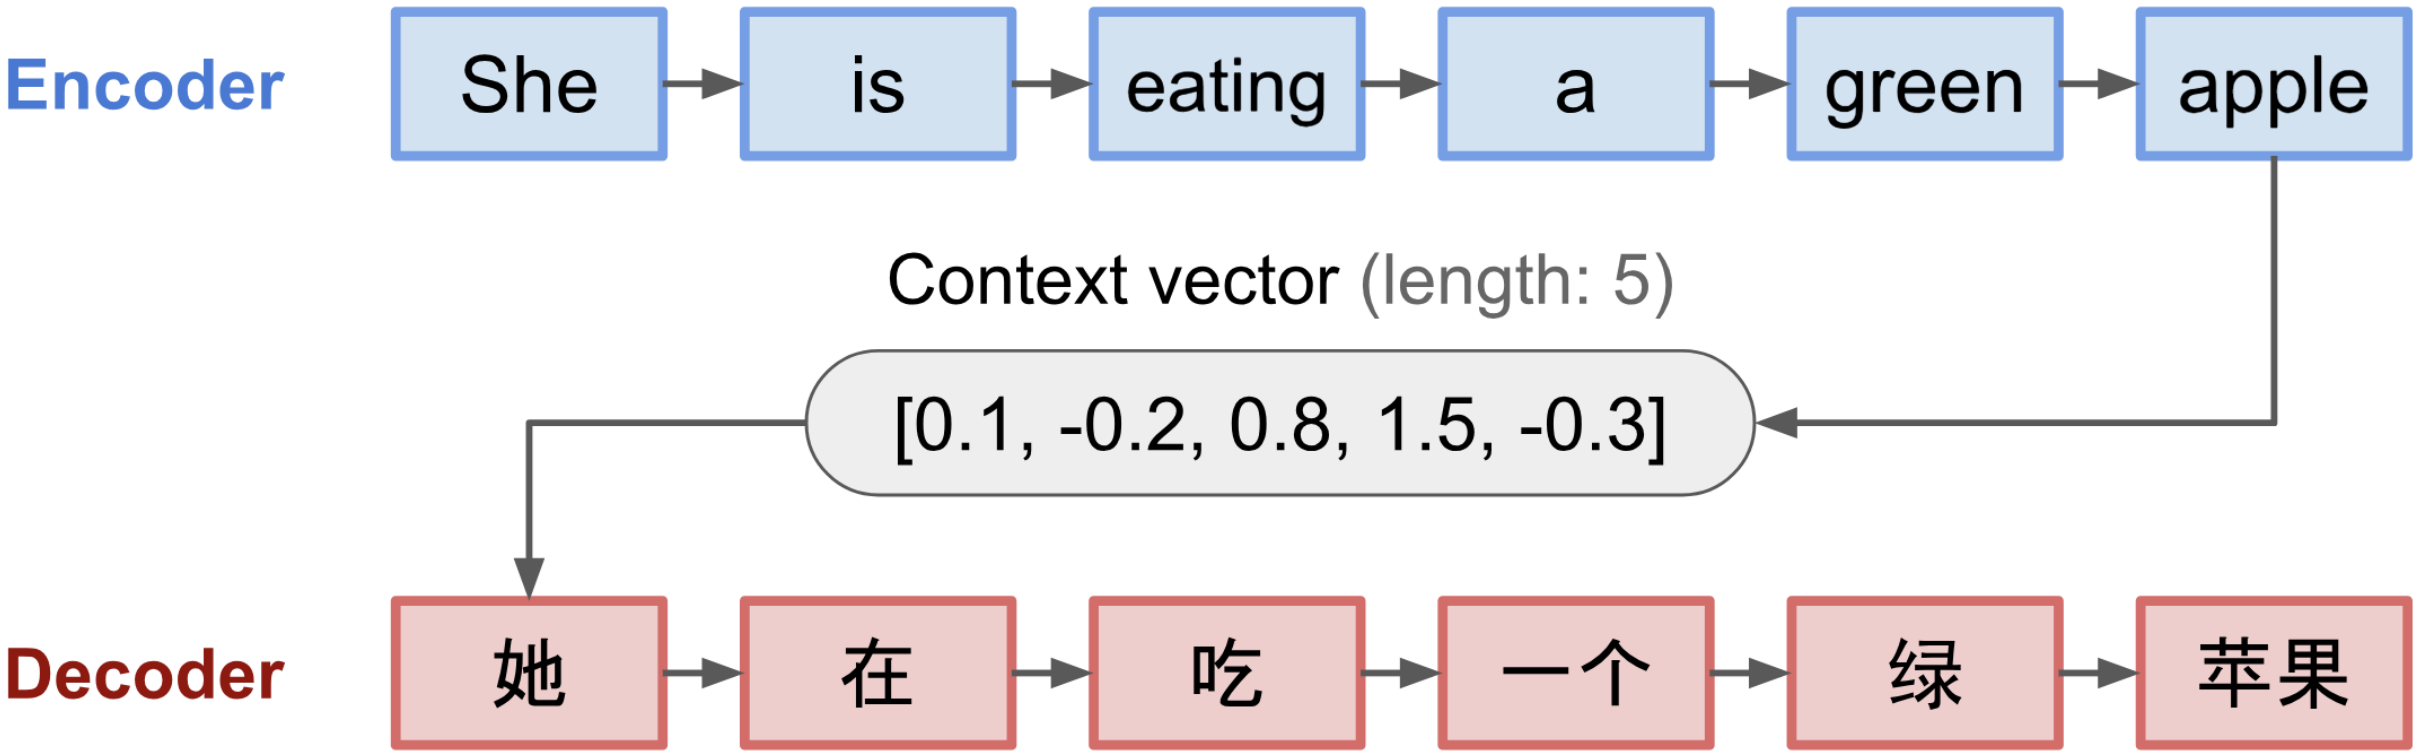

Attention = Vector of Importance Weights

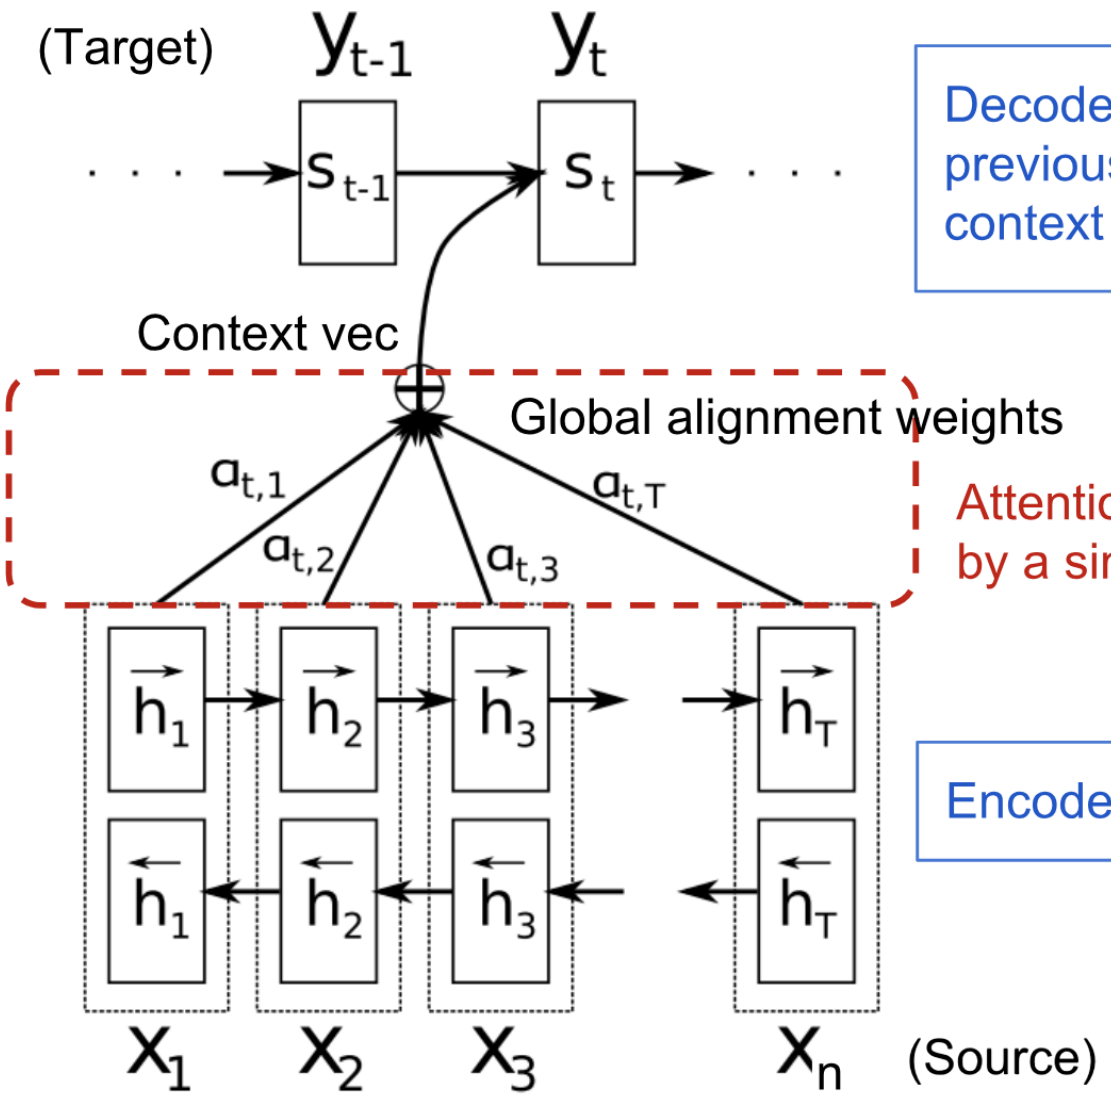

Decoder: RNN with input from previous state + dynamic context vector.

Attention layer: parameterized by a simple feed-forward network

Additive Attention

Encoder: bidirectional RNN

Transformer

- http://jalammar.github.io/illustrated-transformer/

Encoder and Decoder

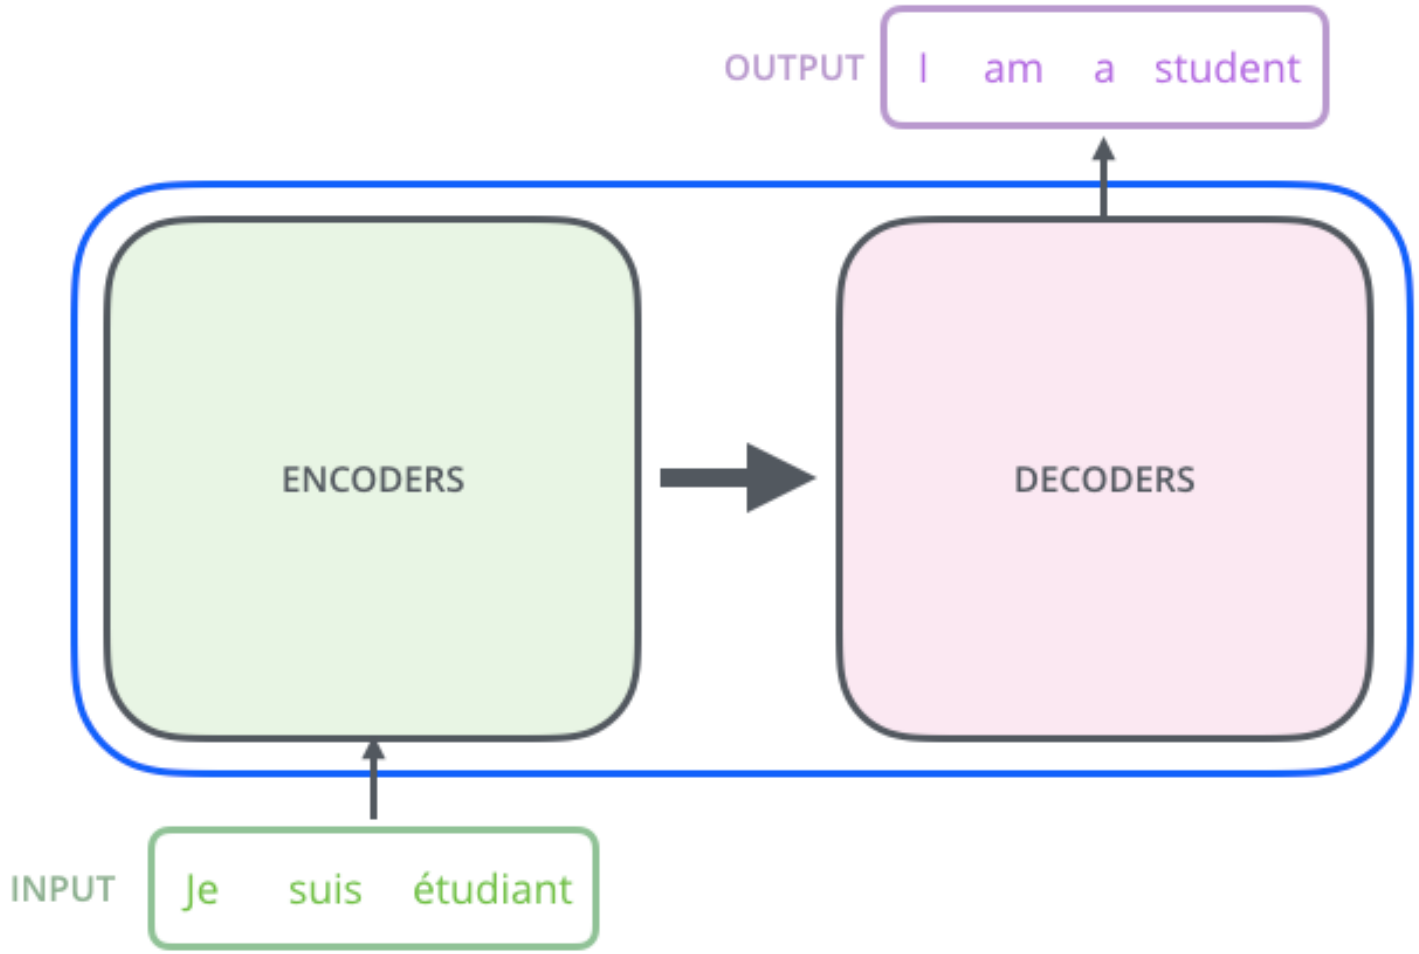

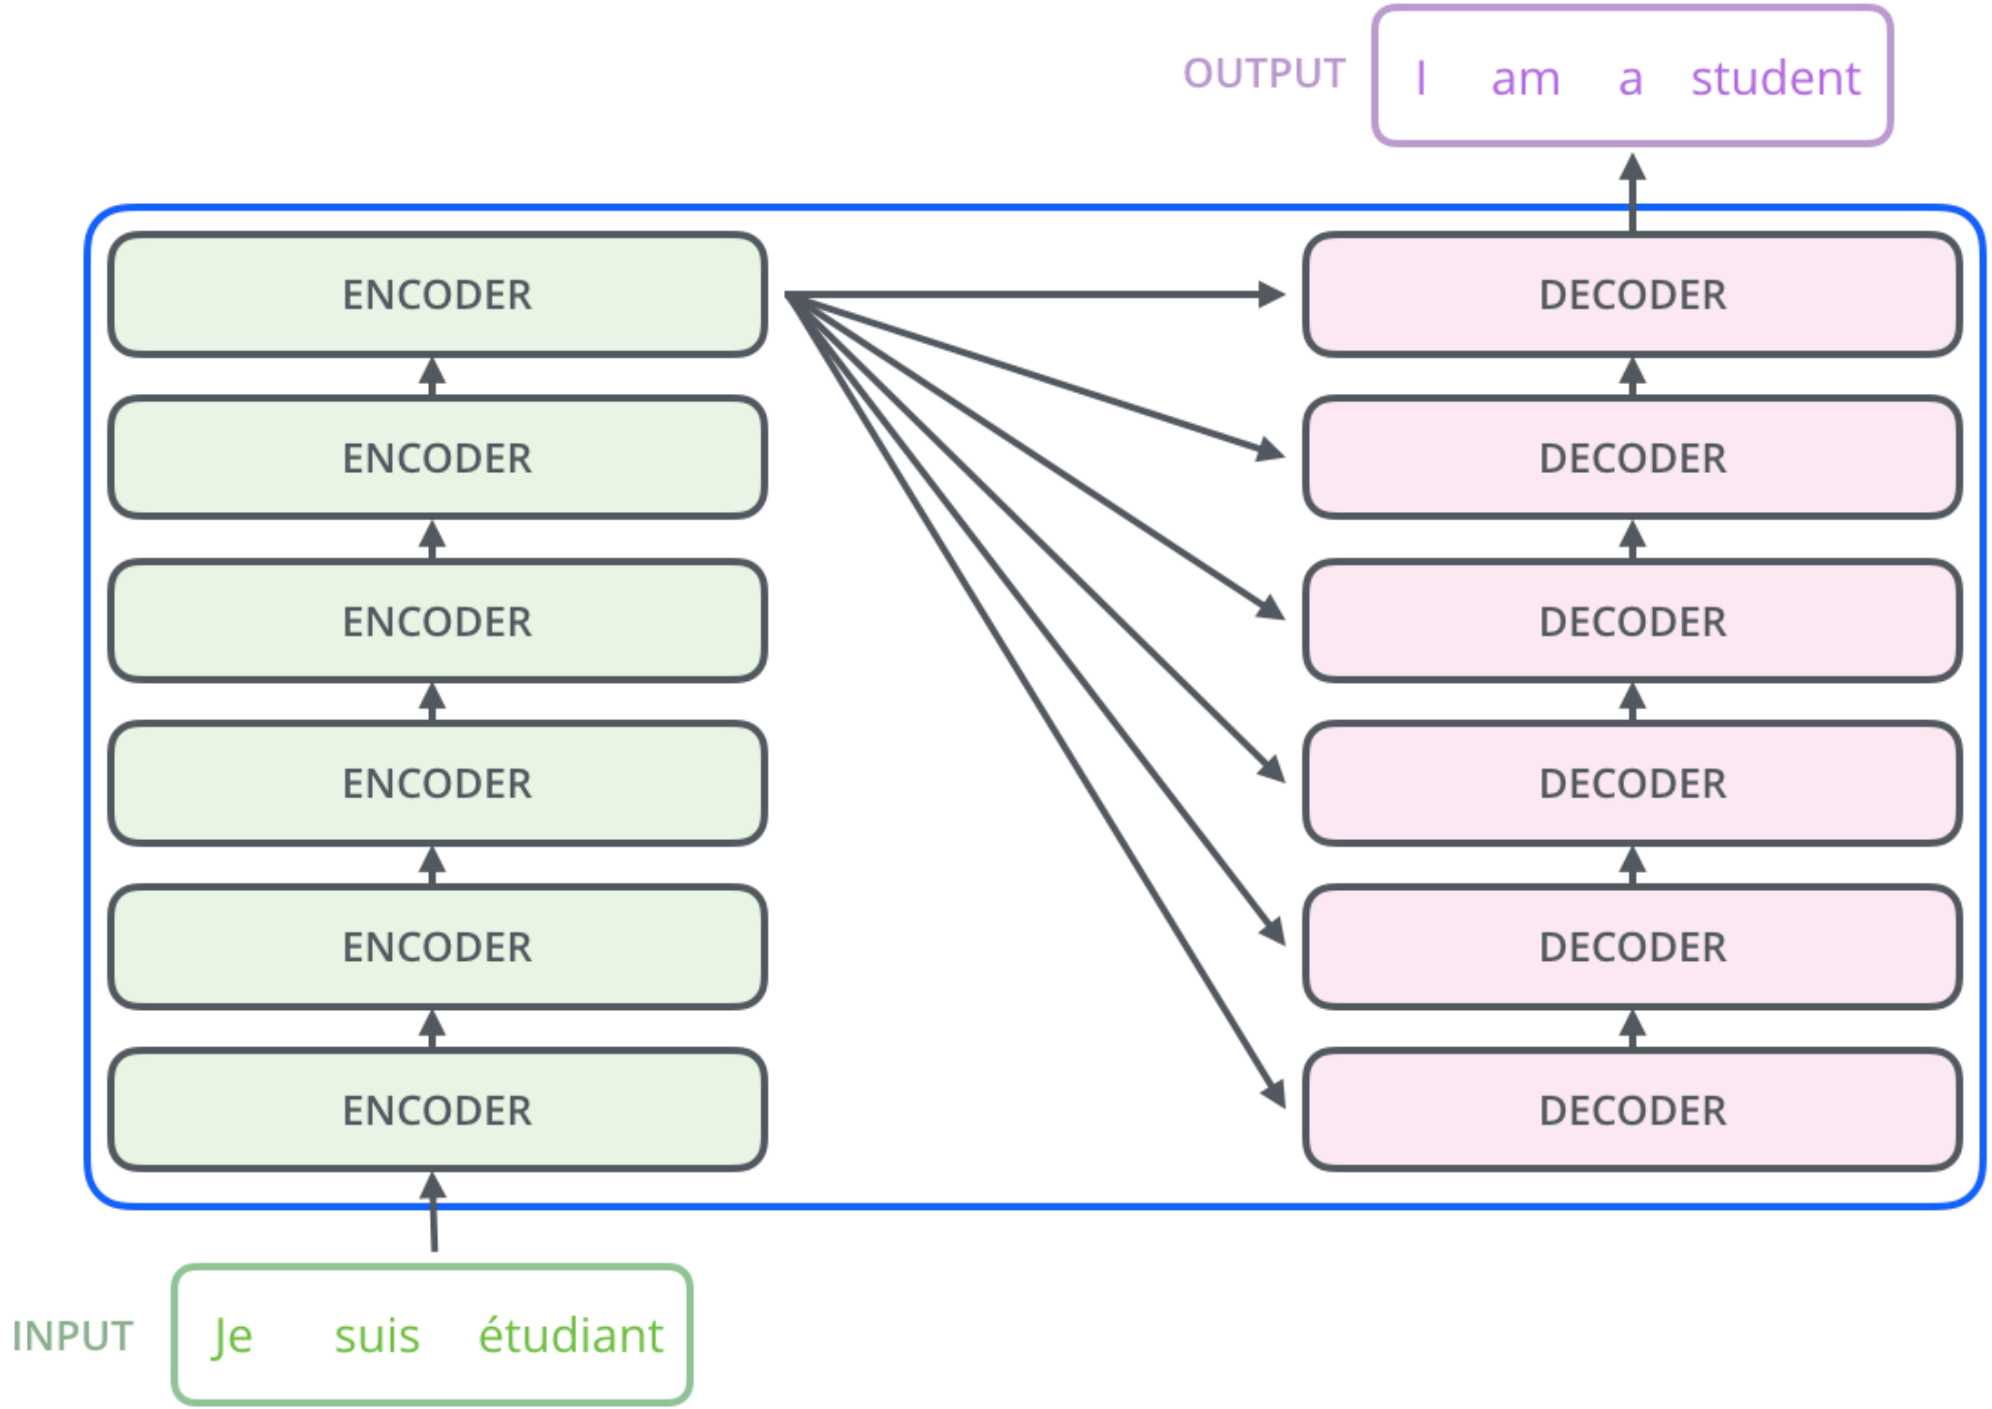
OUTPUT
I am a student
ENCODER
ENCODER
ENCODER
ENCODER
ENCODER
ENCODER
DECODER
DECODER
DECODER
DECODER
DECODER
DECODER
INPUT
Je suis étudiant

Structure of the Encoder and Decoder

- Self-attention
- Encoder-decoder attention

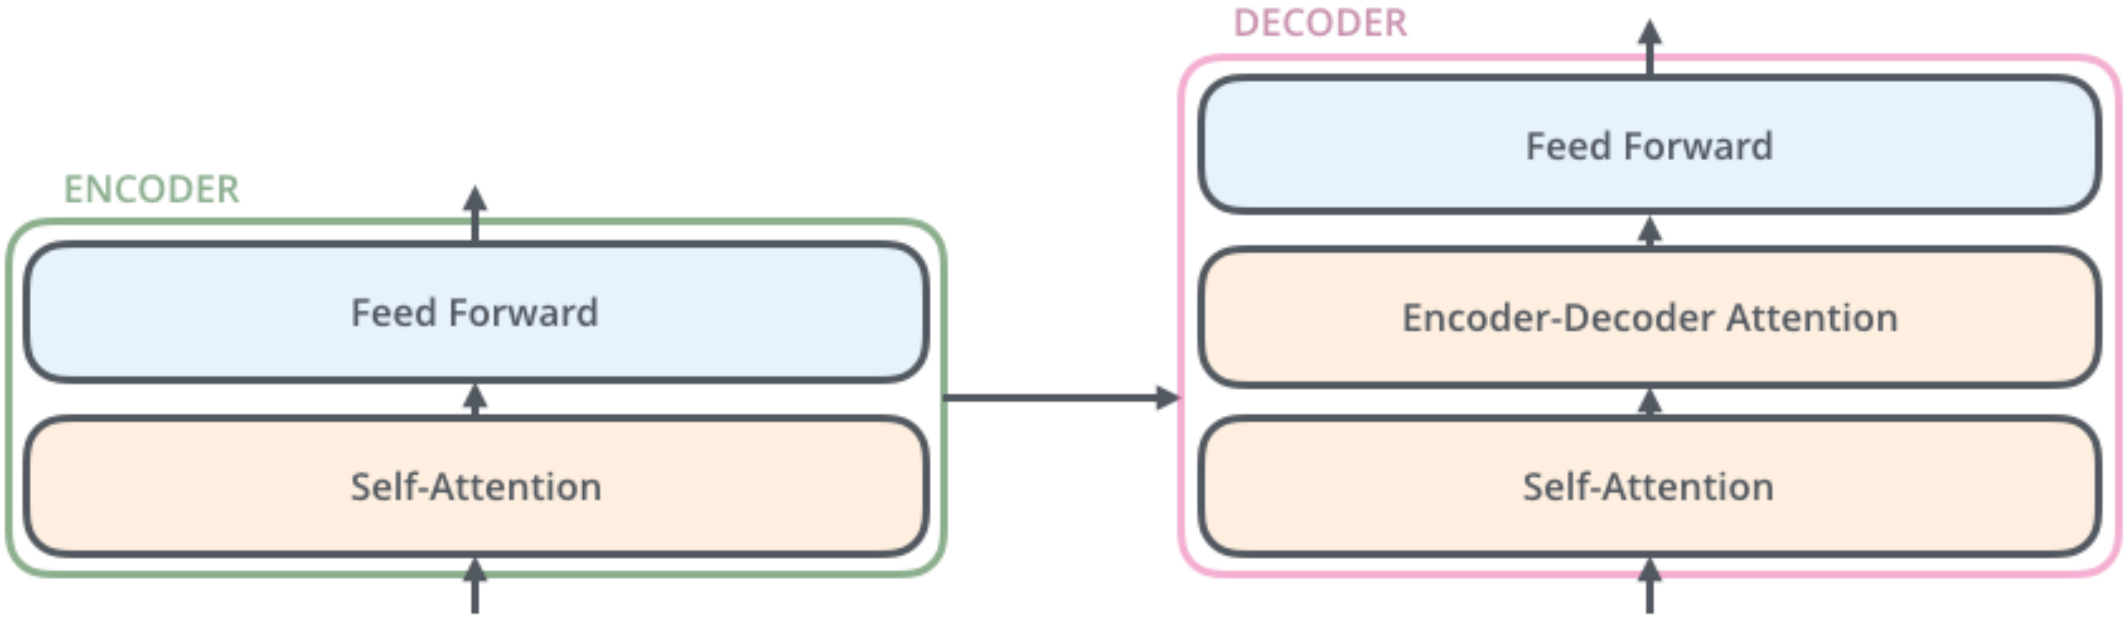

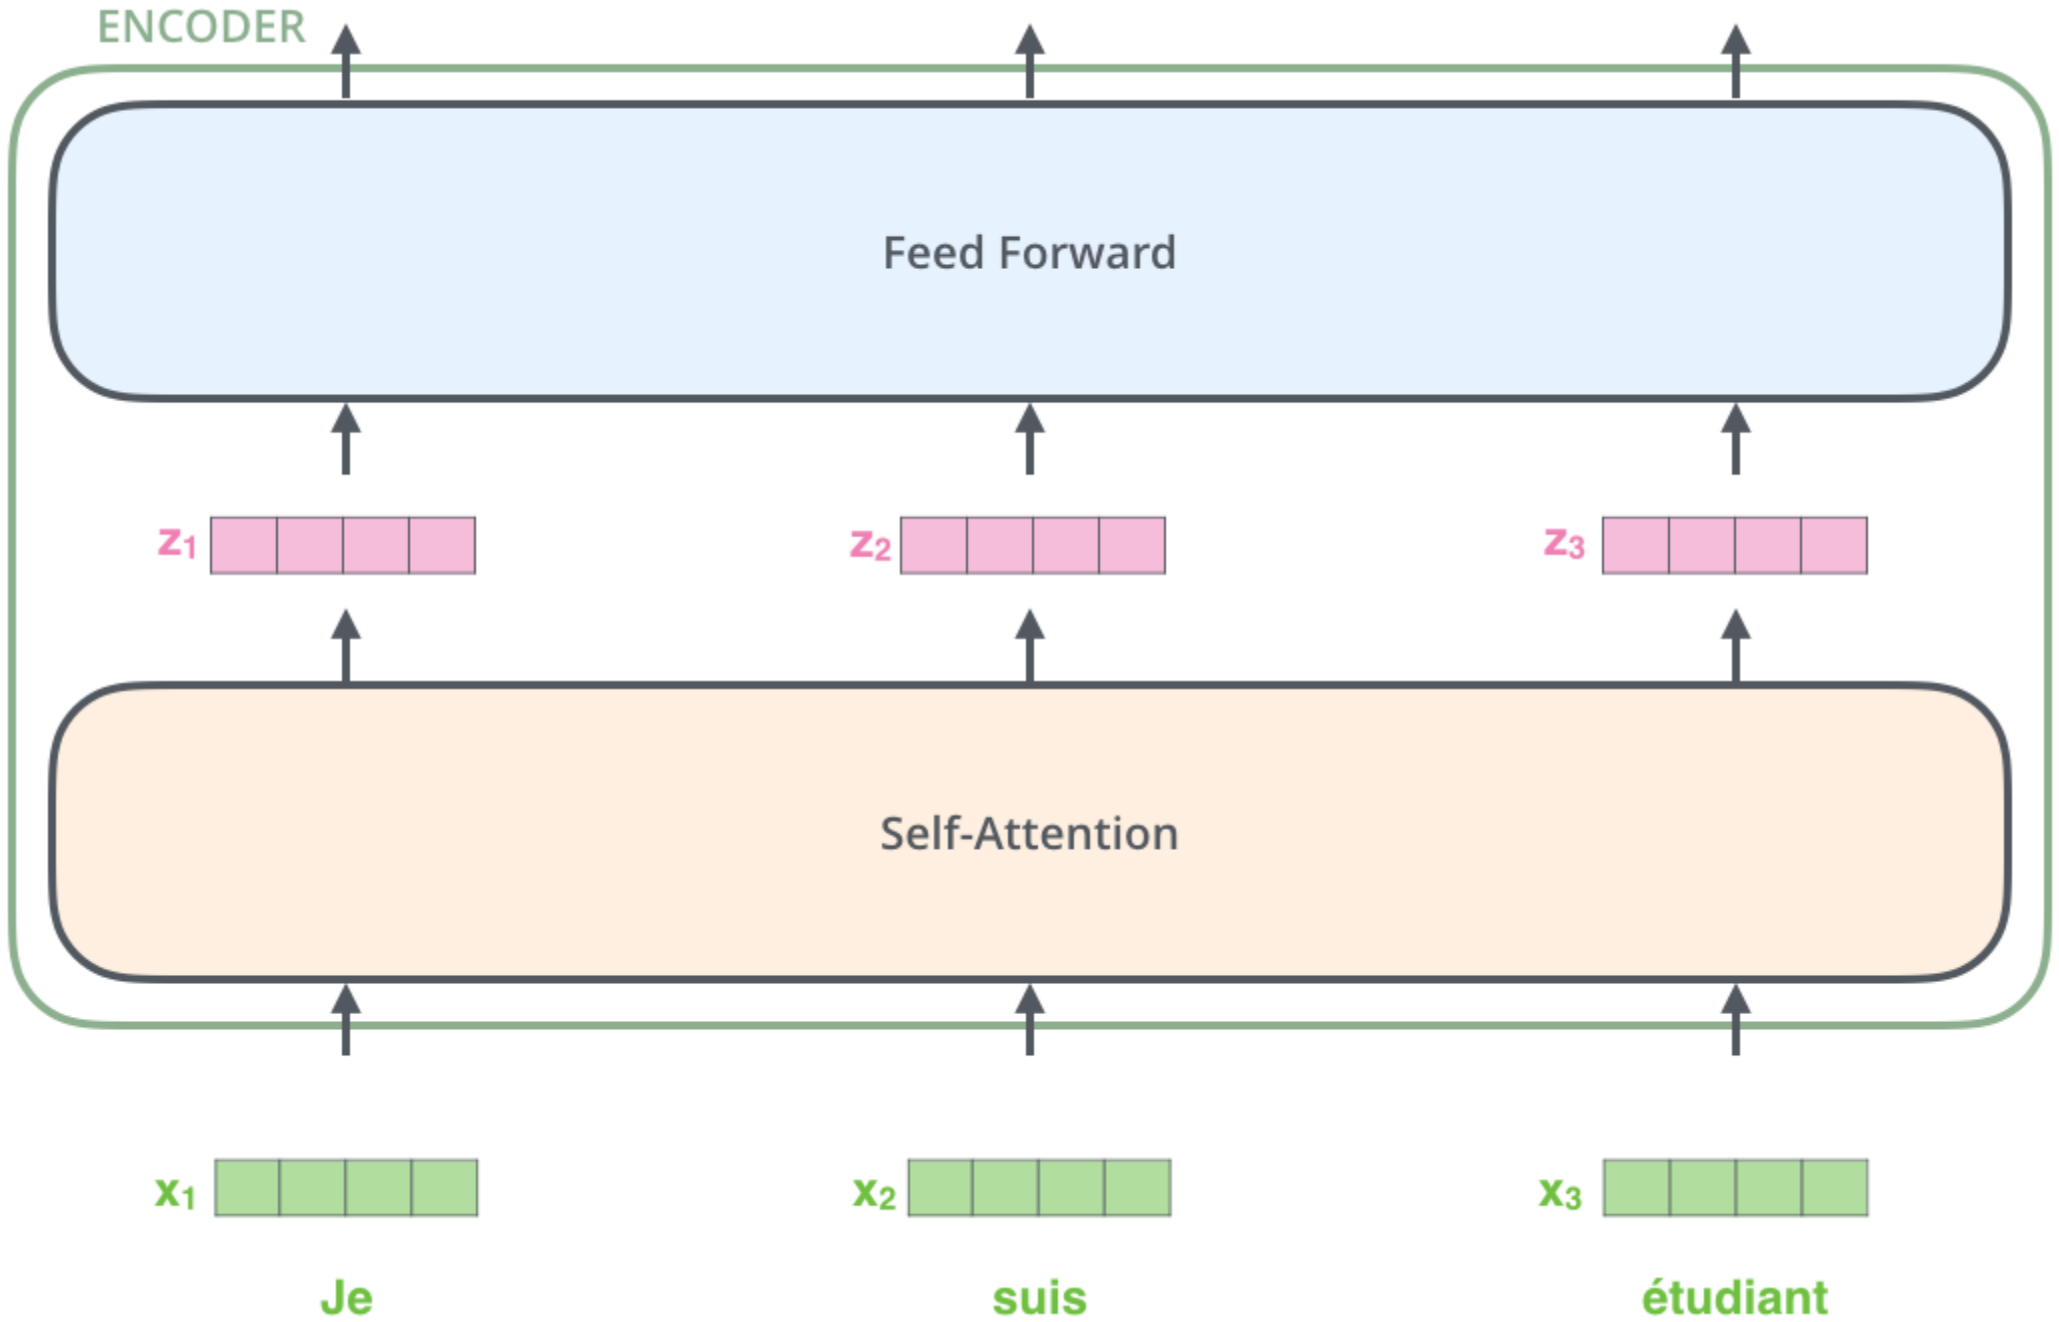
ENCODER
Feed Forward
z1
z2
z3
Self-Attention
x1
x2
x3
Je
suis
étudiant

Tensor2Tensor Notebook

- https://colab.research.google.com/github/tensorflow/tensor2tensor/blob/master/tensor2tensor/notebooks/hello_t2t.ipynb

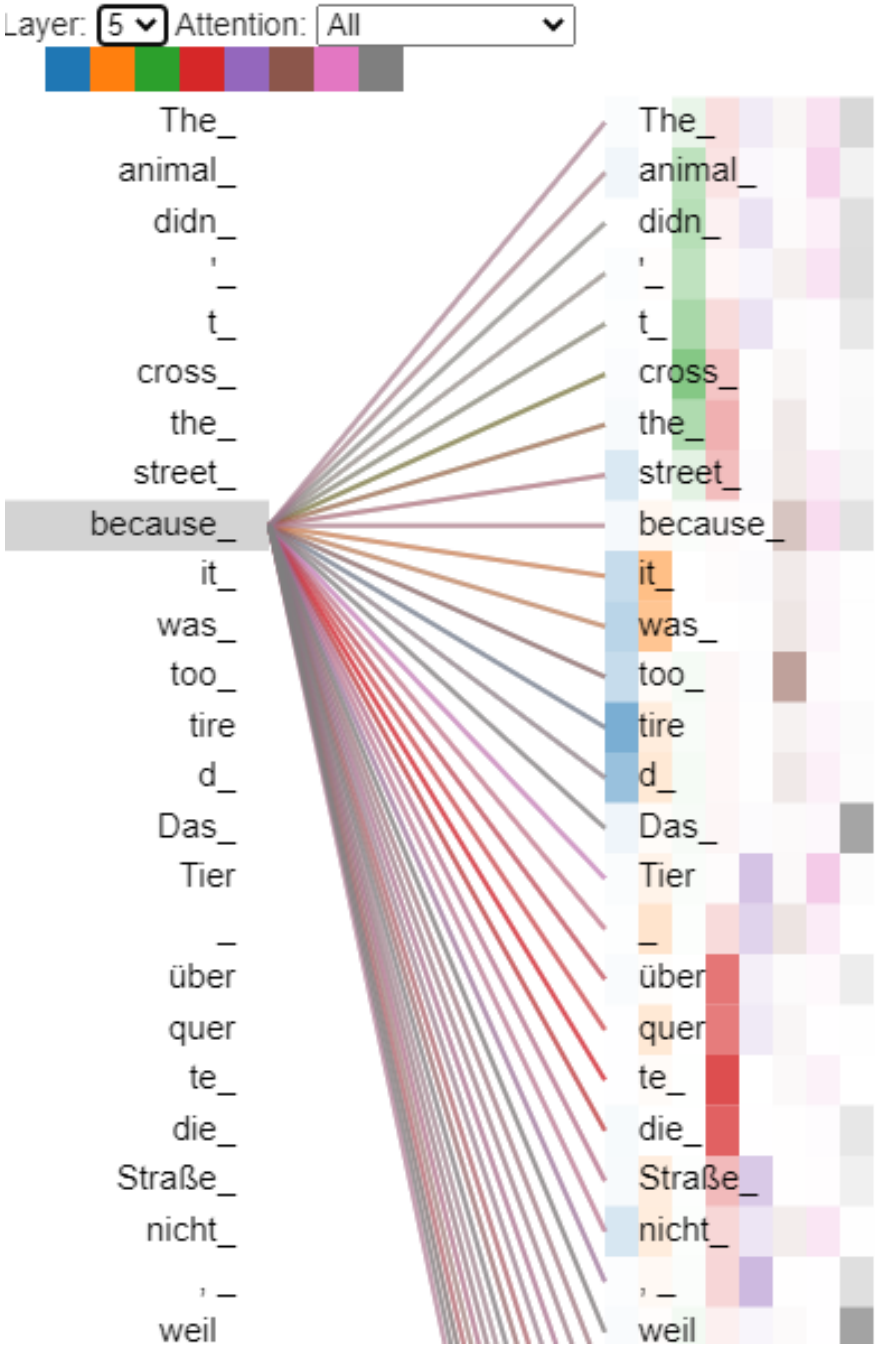

Self-attention (query, key, value)

q: query (to match others)

$$q^i = W^q a^i$$

k: key (to be matched)

$$k^i = W^k a^i$$

v: information to be extracted

$$v^i = W^v a^i$$

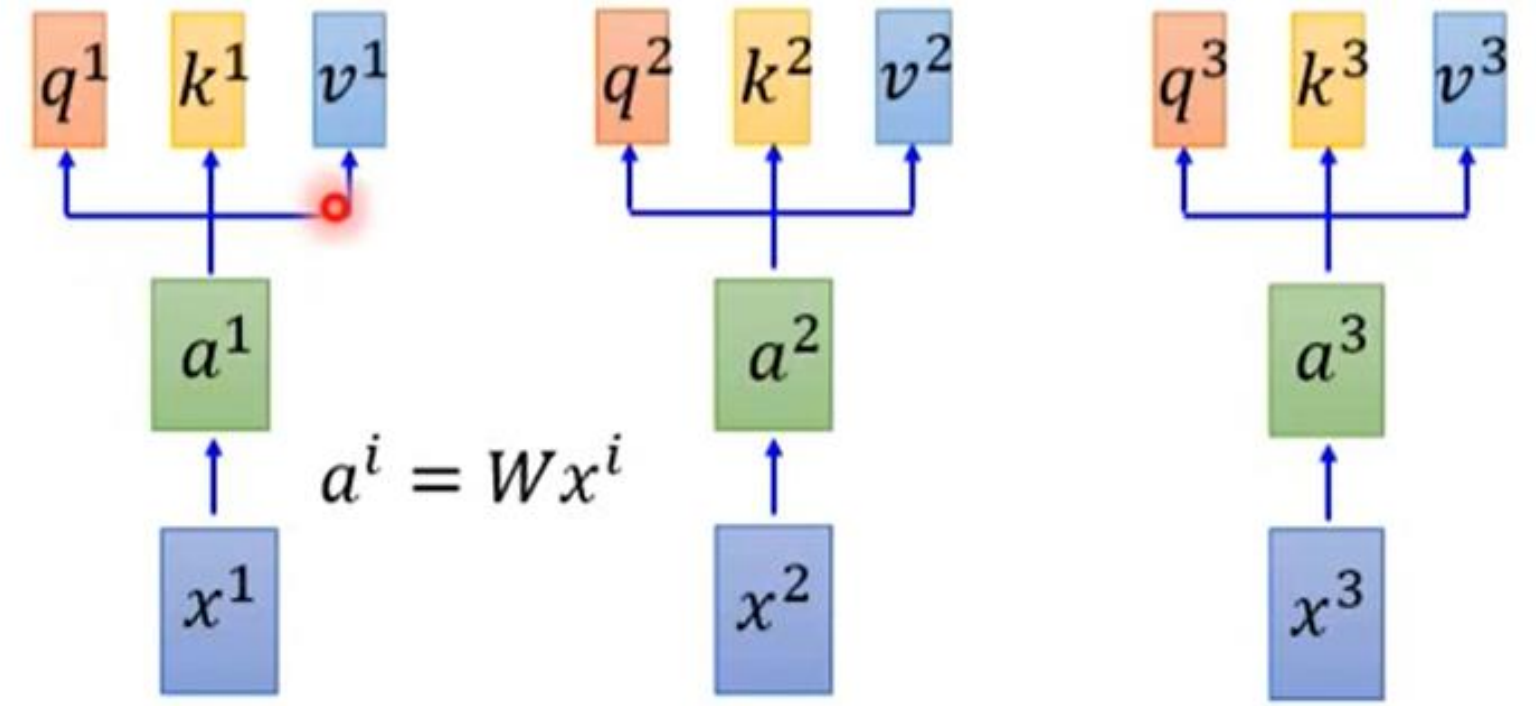

https://www.youtube.com/watch?v=ugWDIIOHtPA&t=1089s

Self-attention

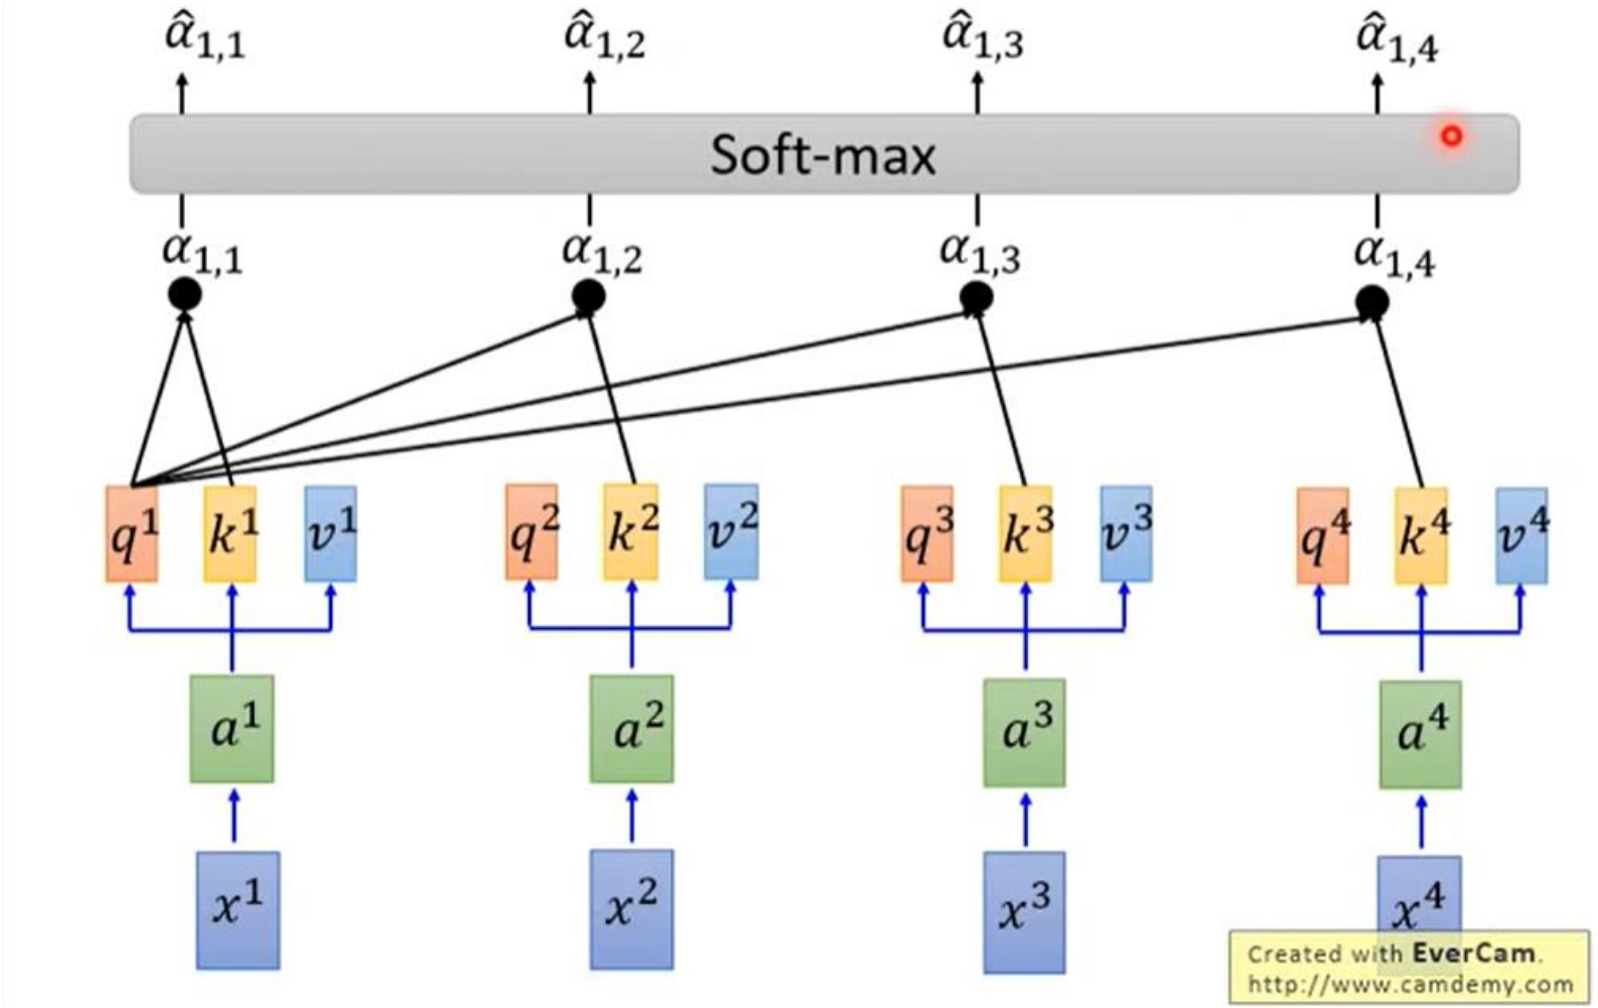

Self-attention

Considering the whole sequence

$$b^1 = \sum_i \hat{\alpha}_{1,i} v^i$$

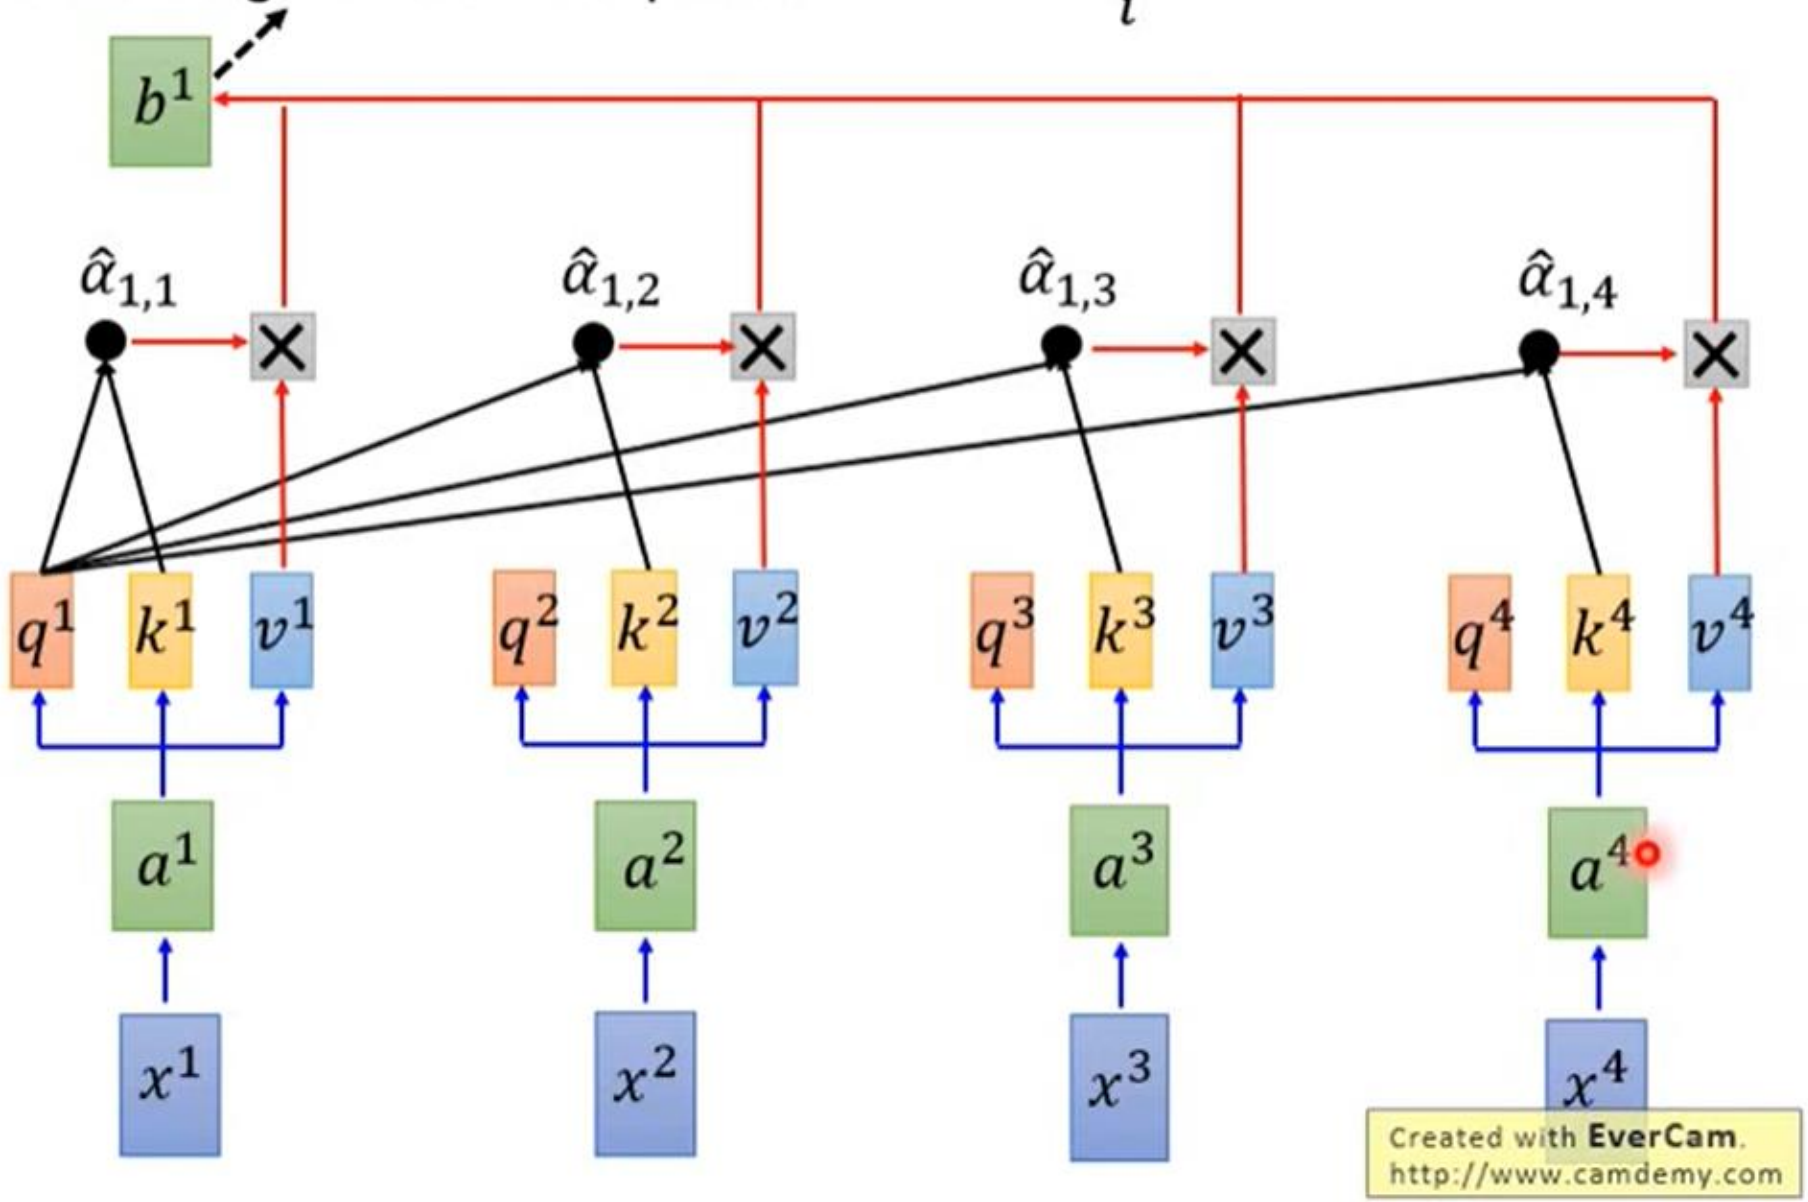

Calculating b^2

$$b^2 = \sum_i \hat{\alpha}_{2,i} v^i$$

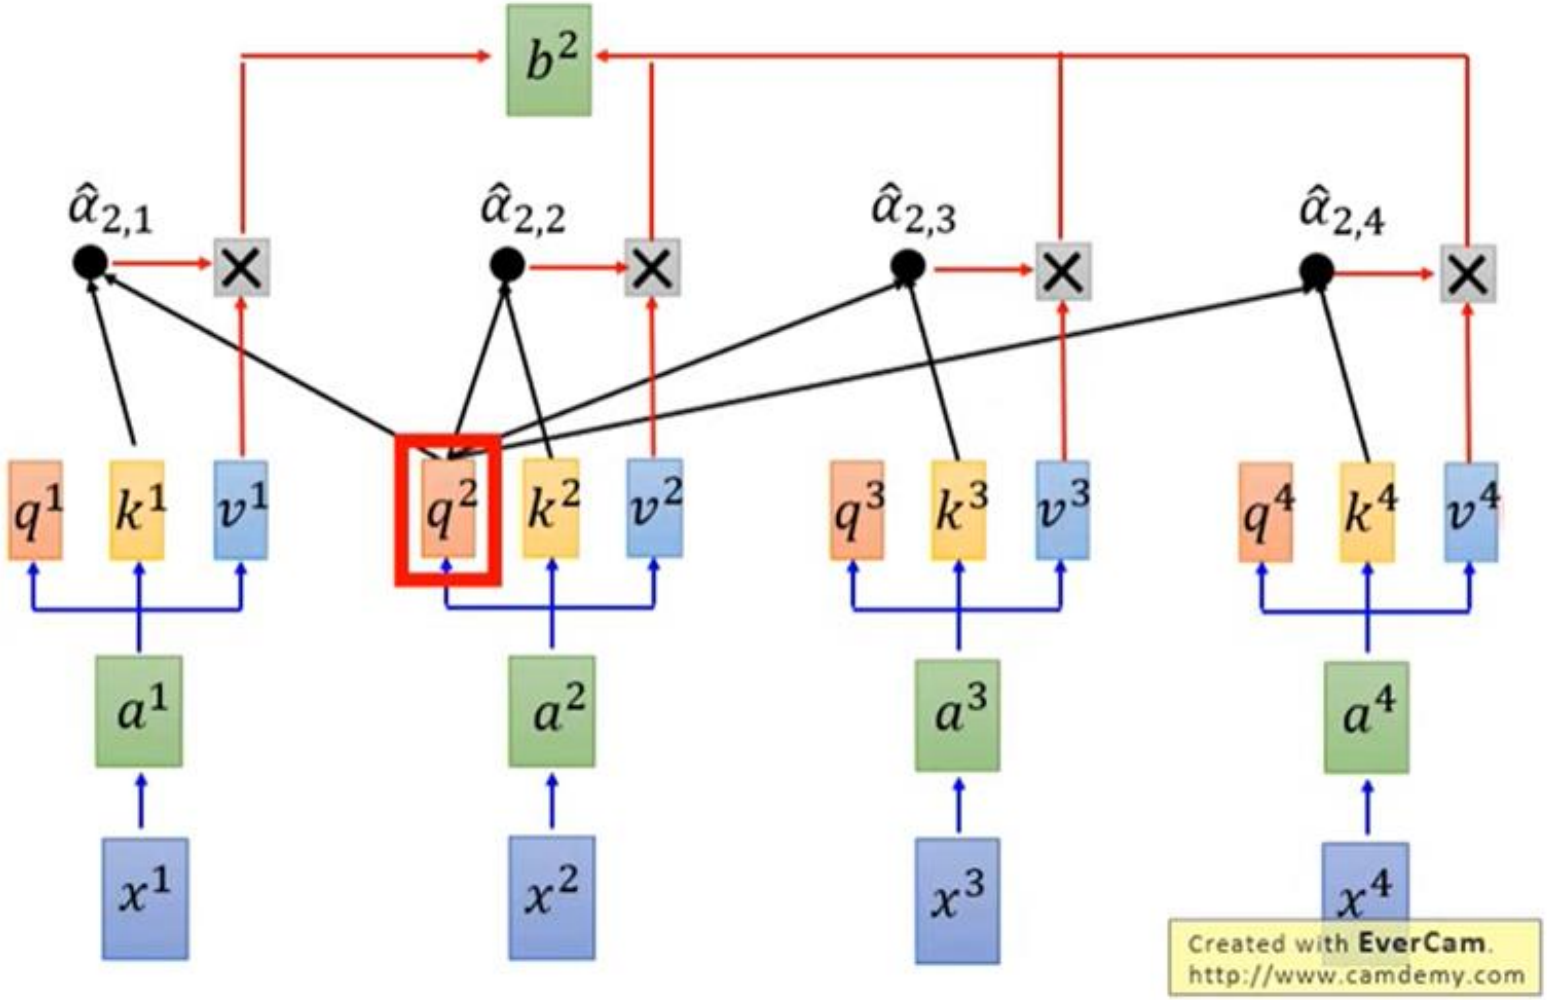

Matrix Mutiplication

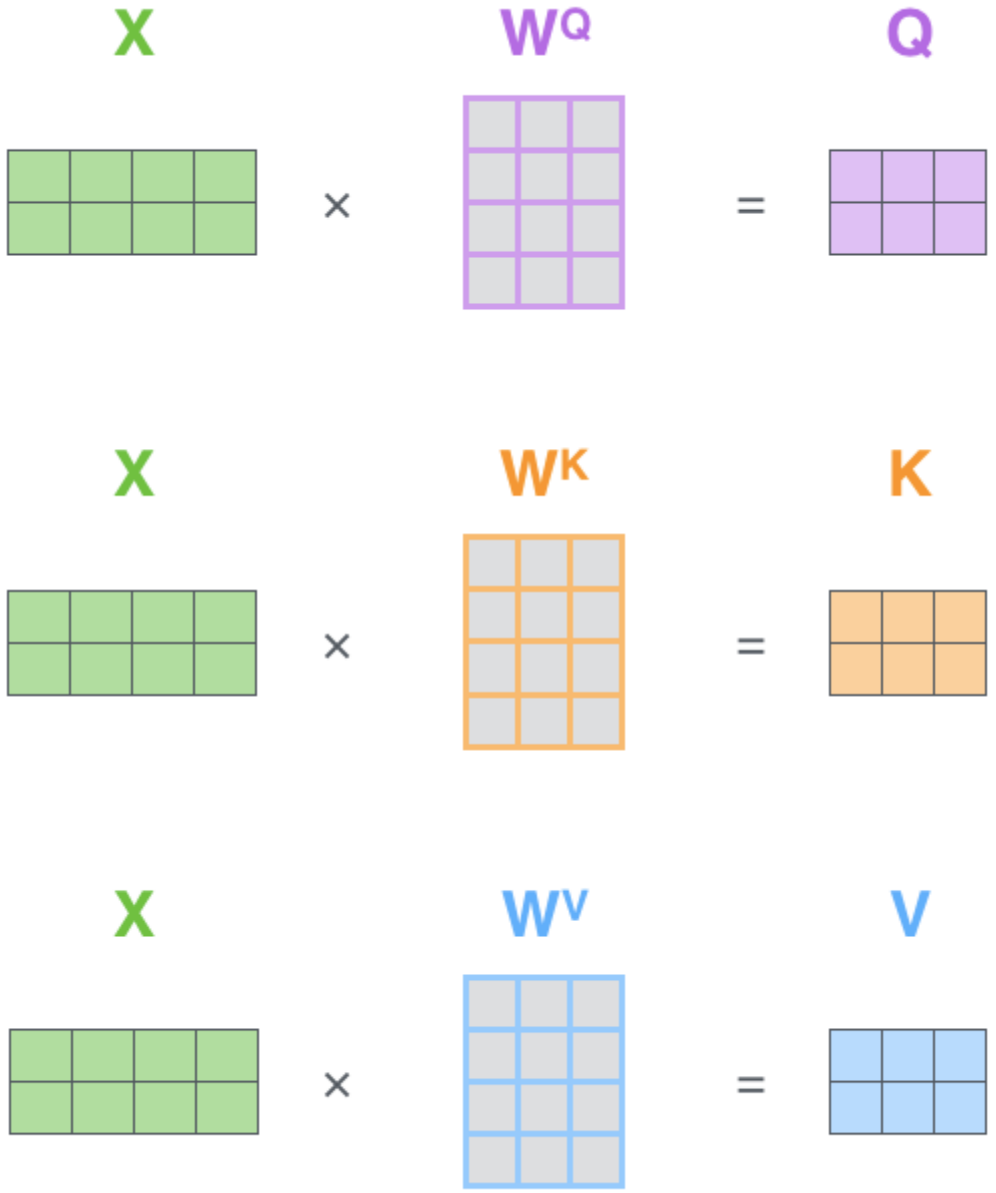

1) This is our input sentence*

2) We embed each word*

3) Split into 8 heads. We multiply X or R with weight matrices

4) Calculate attention using the resulting Q/K/V matrices

5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer

Thinking Machines

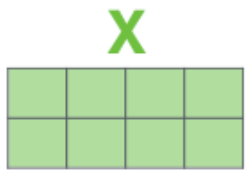

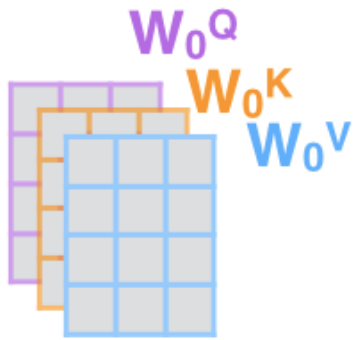

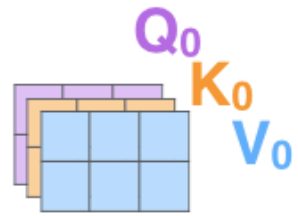

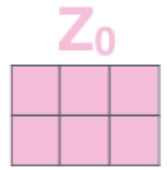

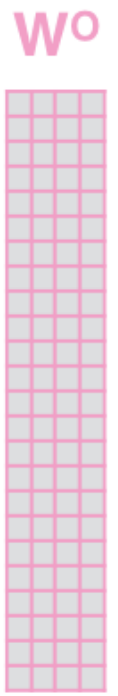

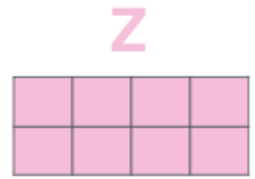

* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one

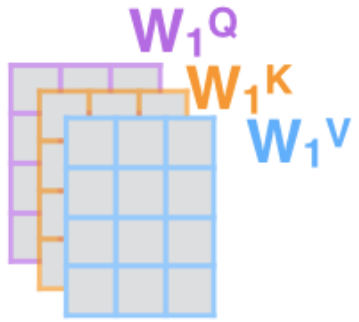

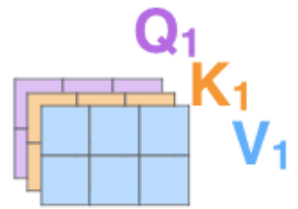

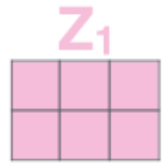

...

...

...

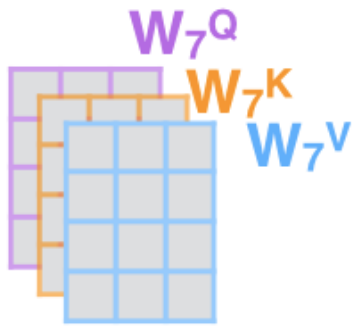

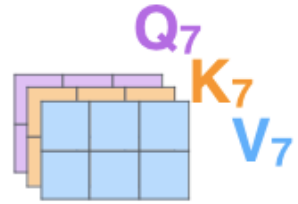

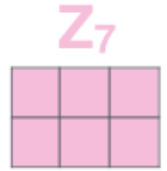

Adding Residual Connections

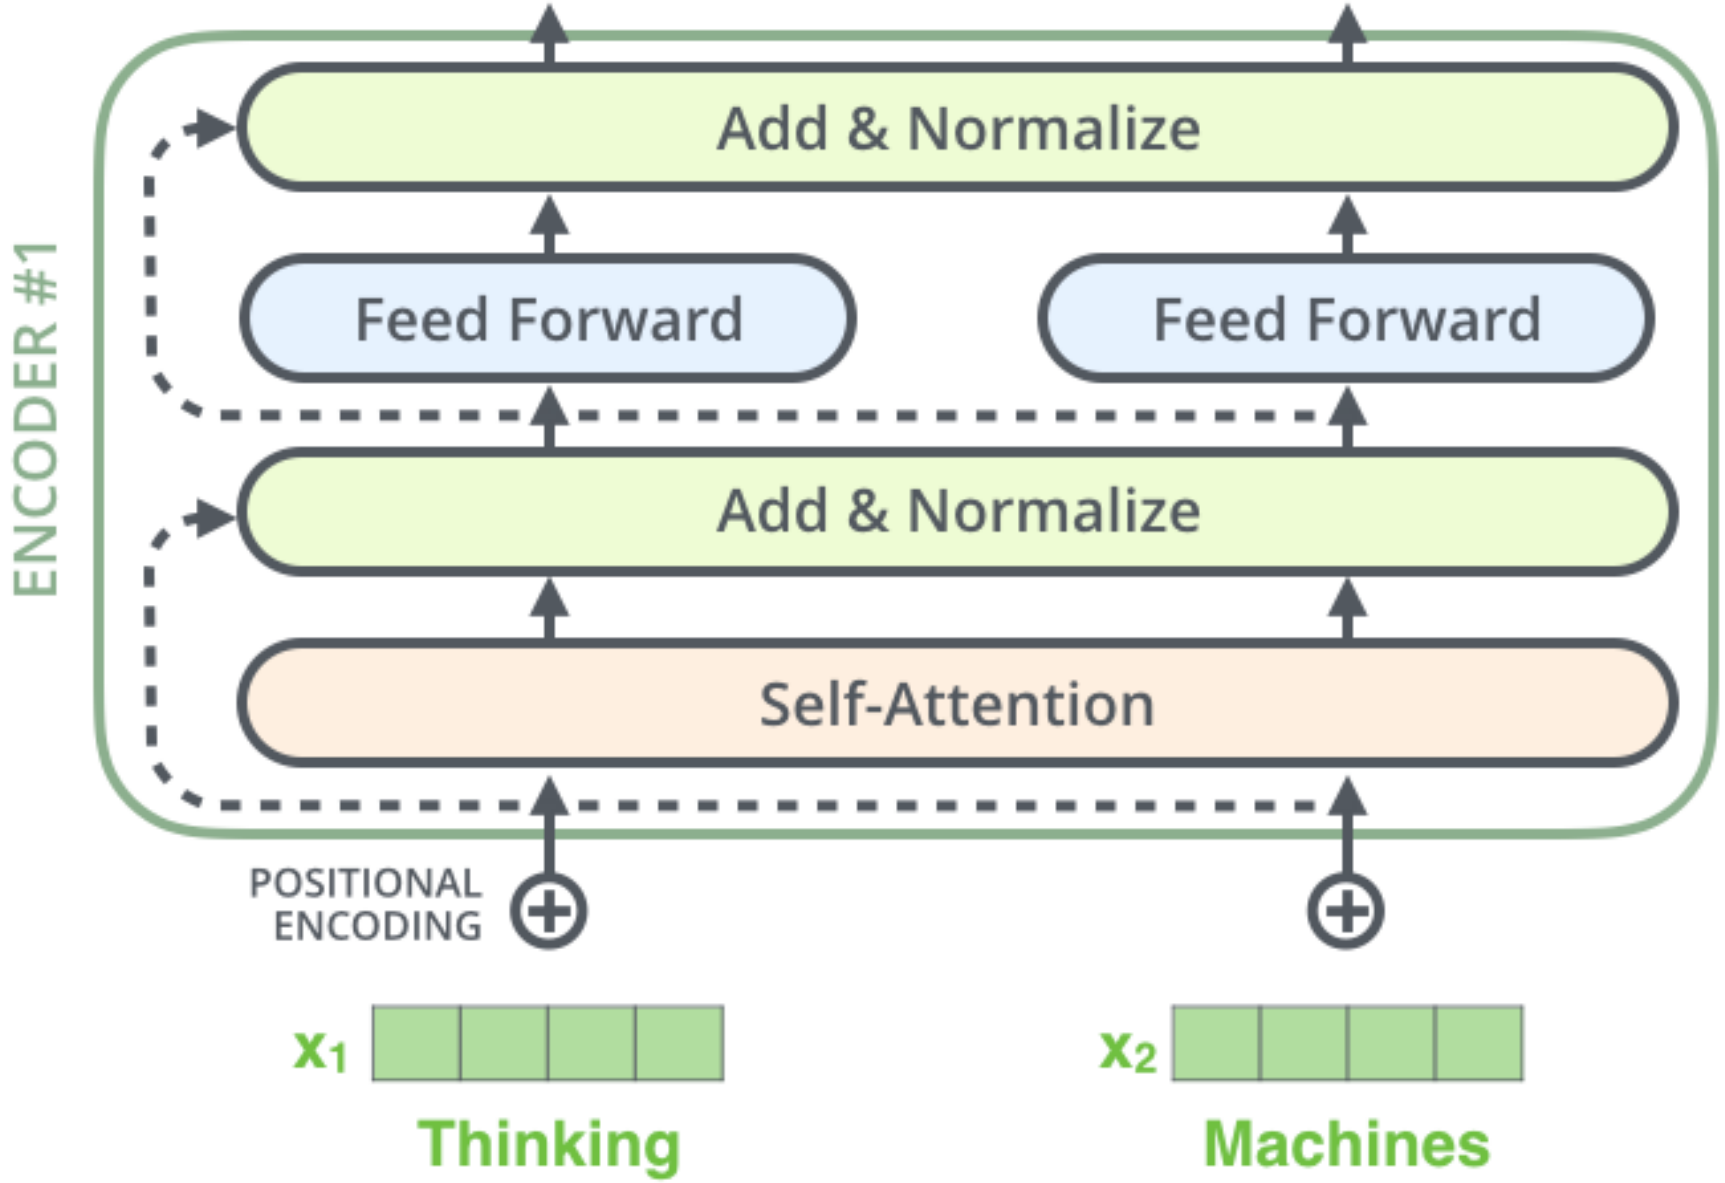

Layer Normalization

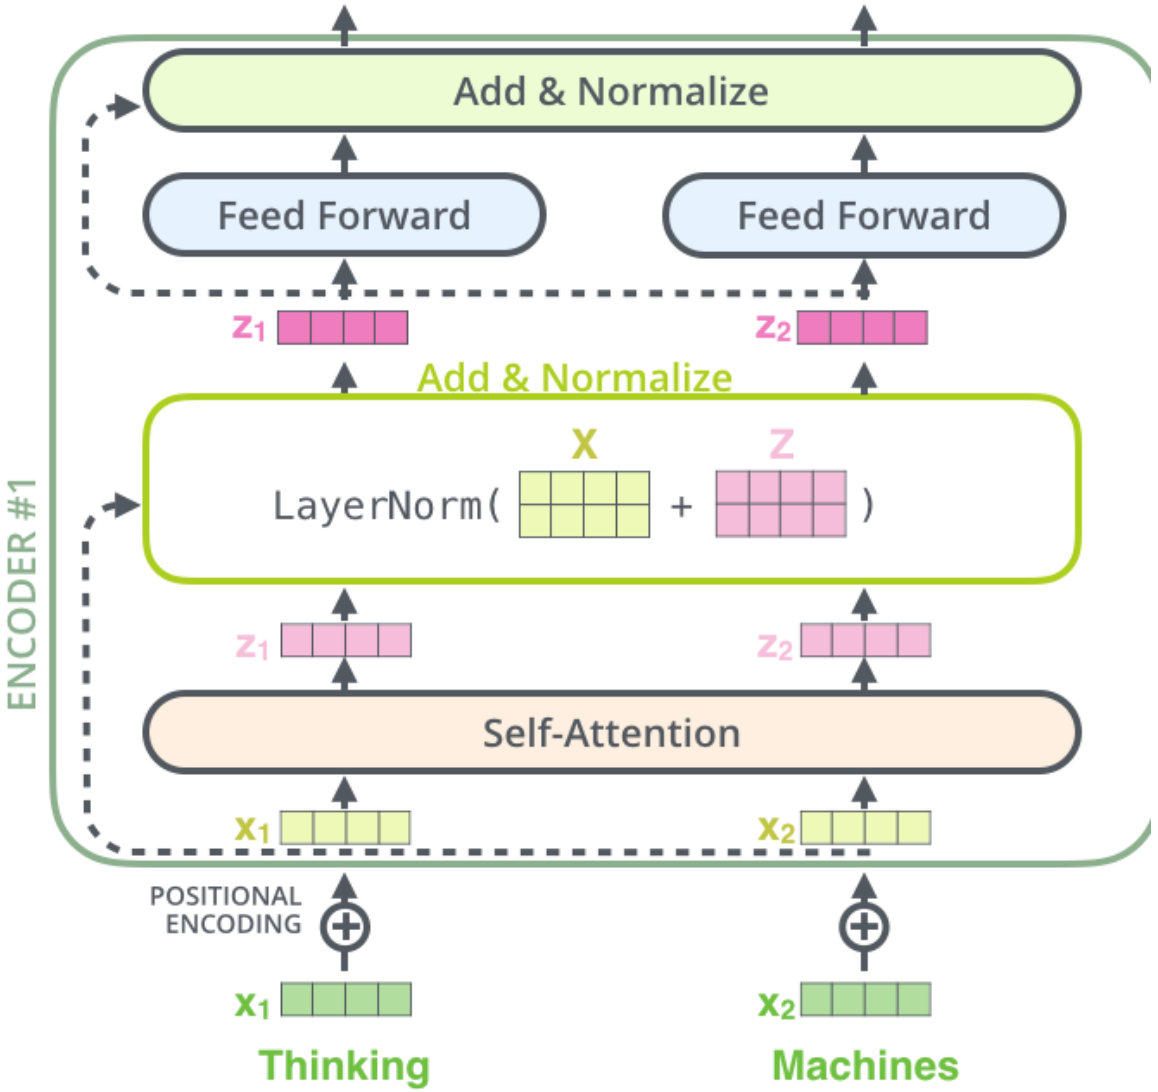

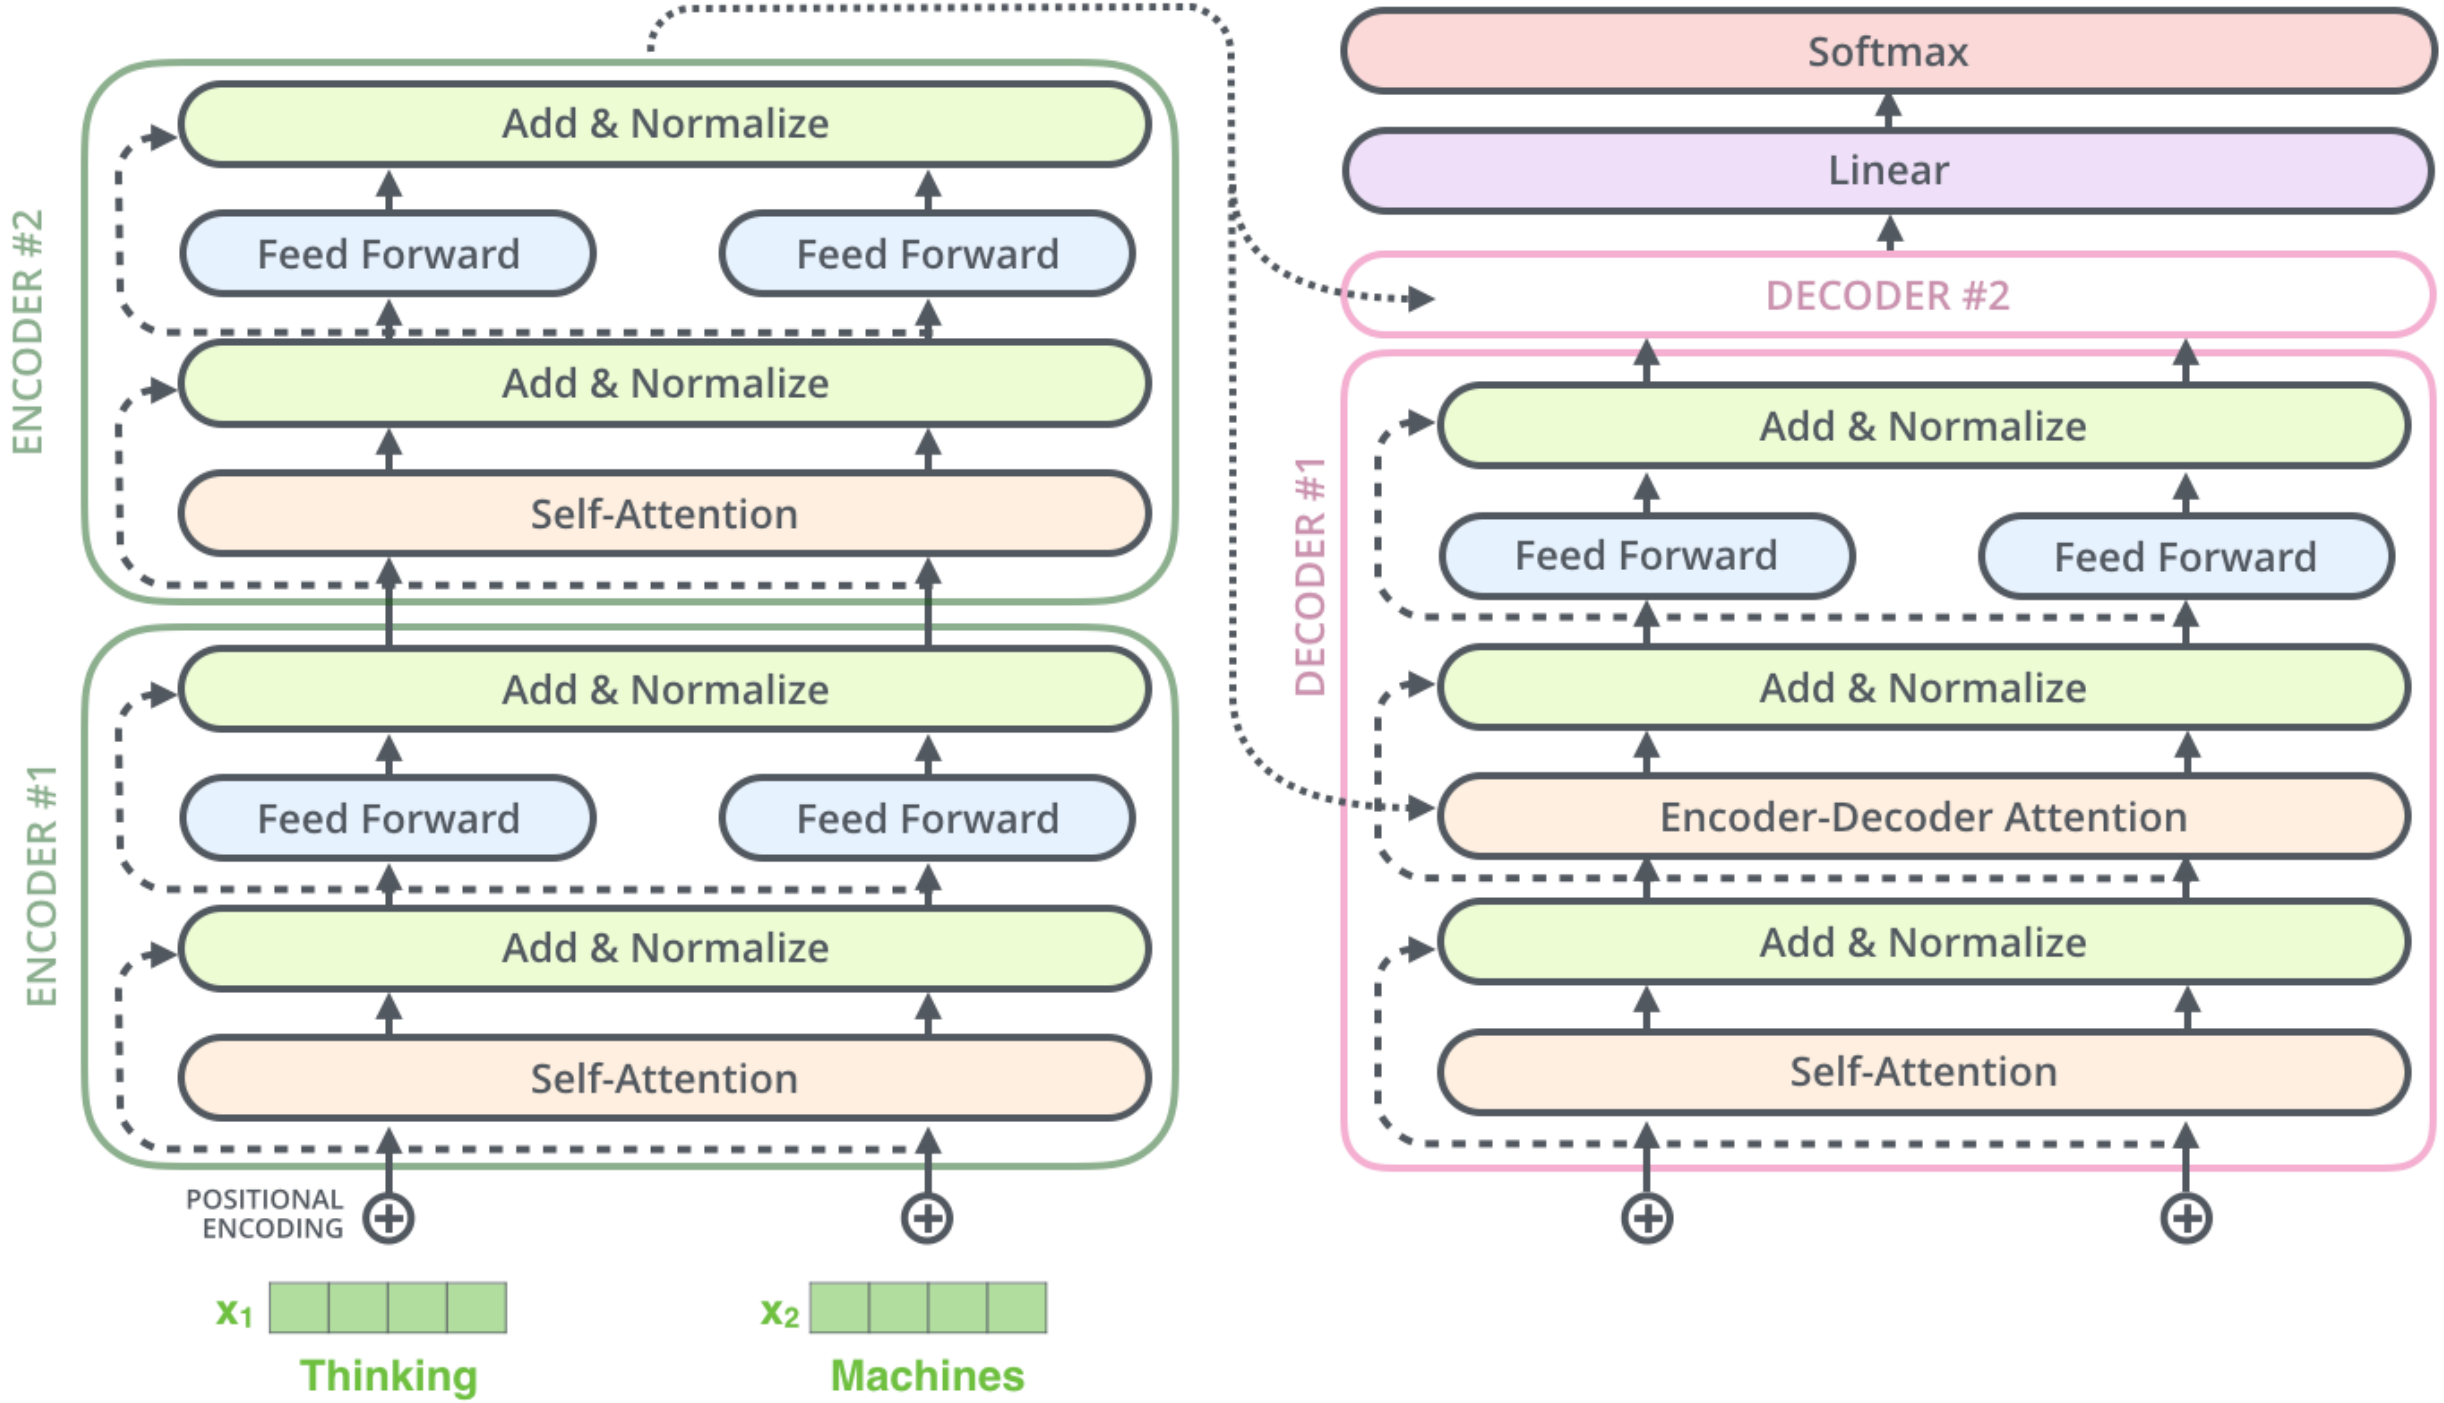
Softmax
Linear
DECODER #2
DECODER #1
Add & Normalize
Feed Forward
Feed Forward
Add & Normalize
Encoder-Decoder Attention
Add & Normalize
Self-Attention
ENCODER #2
Add & Normalize
Feed Forward
Feed Forward
Add & Normalize
Self-Attention
ENCODER #1
Add & Normalize
Feed Forward
Feed Forward
Add & Normalize
Self-Attention
POSITIONAL ENCODING
x1
x2
Thinking
Machines

References

1. https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html
2. http://jalammar.github.io/illustrated-transformer/
3. Hong-Yi Lee, Transformer, 2019
   https://www.youtube.com/watch?v=ugWDIIOHtPA