


# Advanced Keras API

Prof. Kuan-Ting Lai

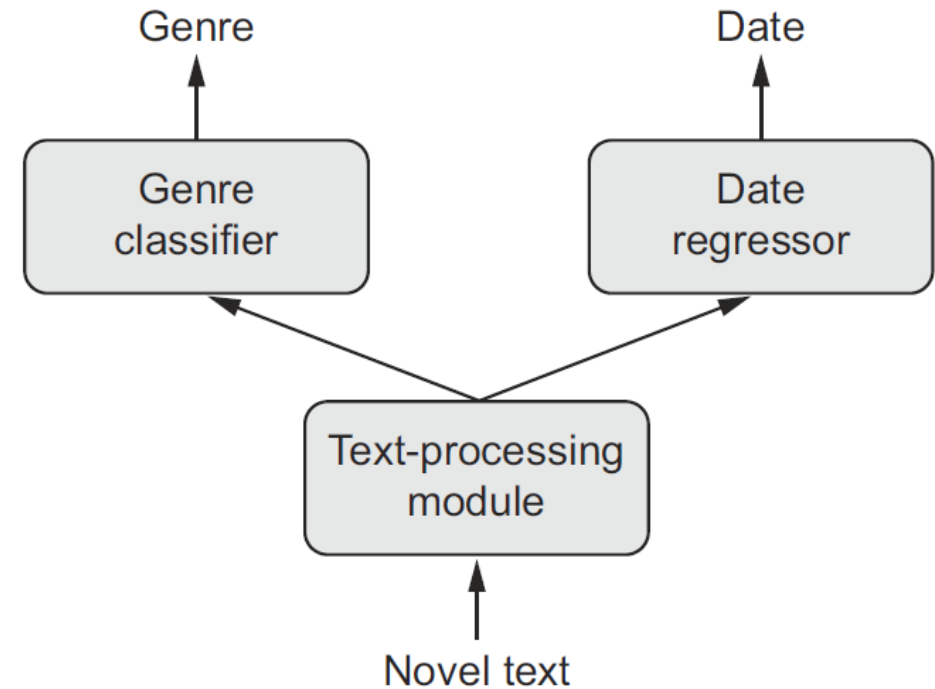
2019/12/24

# Going Beyond Sequential Model

## Multi-input Model




## Multi-output Model




# Inception Module

- Christian Szegedy et al., “Going Deeper with Convolutions,” *CVPR*, 2014.  
<https://arxiv.org/abs/1409.4842>



# Residual Connection

- Kaiming He et al., “Deep Residual Learning for Image Recognition,” CVPR (2015), <https://arxiv.org/abs/1512.03385>



# Functional API


```
from keras import Input, layers

input_tensor = Input(shape=(32,))

dense = layers.Dense(32, activation='relu')

output_tensor = dense(input_tensor)
```

A layer is a function



A layer may be called on a tensor, and it returns a tensor

# Functional API vs. Sequential Model

- Create a Model object using only an input tensor and an output tensor

```
input_tensor = Input(shape=(64,))  
x = layers.Dense(32, activation='relu')(input_tensor)  
x = layers.Dense(32, activation='relu')(x)  
output_tensor = layers.Dense(10, activation='softmax')(x)  
model = Model(input_tensor, output_tensor)
```



# Functional API vs. Sequential Model

- Create a Model object using only an input tensor and an output tensor

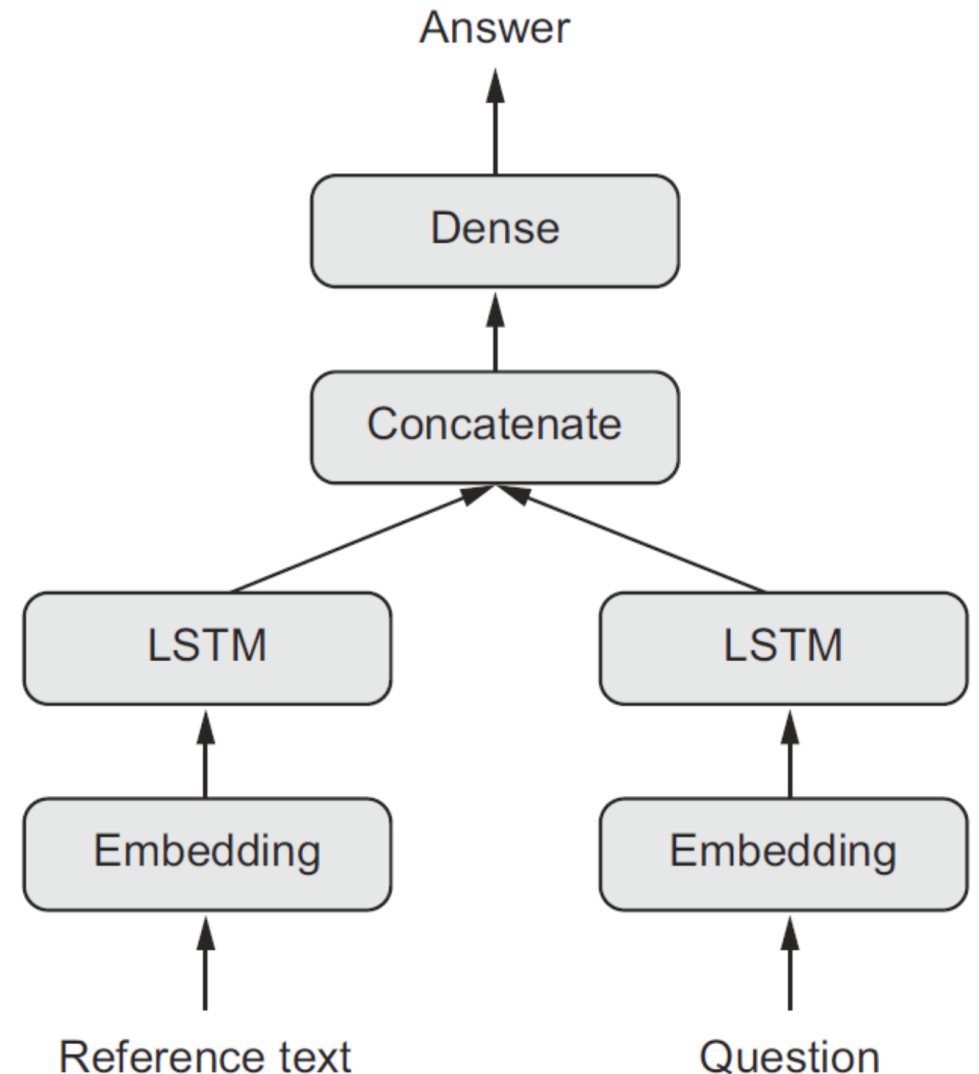
```
input_tensor = Input(shape=(64,))  
x = layers.Dense(32, activation='relu')(input_tensor)  
x = layers.Dense(32, activation='relu')(x)  
output_tensor = layers.Dense(10, activation='softmax')(x)  
model = Model(input_tensor, output_tensor)
```



```
seq_model = Sequential()  
seq_model.add(layers.Dense(32, activation='relu',  
                           input_shape=(64,)))  
seq_model.add(layers.Dense(32, activation='relu'))  
seq_model.add(layers.Dense(10, activation='softmax'))
```

# Question-answering Model

- Two inputs:
  1. A natural-language question
  2. Reference text snippet (such as a news article)
- One output: answer
  - One-word answer obtained via a SoftMax over some predefined vocabulary





```
text_vocabulary_size = 10000
question_vocabulary_size = 10000
answer_vocabulary_size = 500

### Reference text ###
text_input = Input(shape=(None,), dtype='int32', name='text')
embedded_text = layers.Embedding(64, text_vocabulary_size)(text_input)
encoded_text = layers.LSTM(32)(embedded_text)

### Question ###
question_input = Input(shape=(None,), dtype='int32', name='question')
embedded_question = layers.Embedding(32, question_vocabulary_size)(question_input)
encoded_question = layers.LSTM(16)(embedded_question)

### Concatenate ###
concatenated = layers.concatenate([encoded_text, encoded_question], axis=-1)
answer = layers.Dense(answer_vocabulary_size, activation='softmax')(concatenated)

model = Model([text_input, question_input], answer)
model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['acc'])
```

# Train Two-input Models

- Training data can be array or dictionary

```
text = np.random.randint(1, text_vocabulary_size, size=(num_samples, max_length))
question = np.random.randint(1, question_vocabulary_size, size=(num_samples,
                                                                max_length))
answers = np.random.randint(0, 1, size=(num_samples, answer_vocabulary_size))
```

```
### Option 1 ###
```

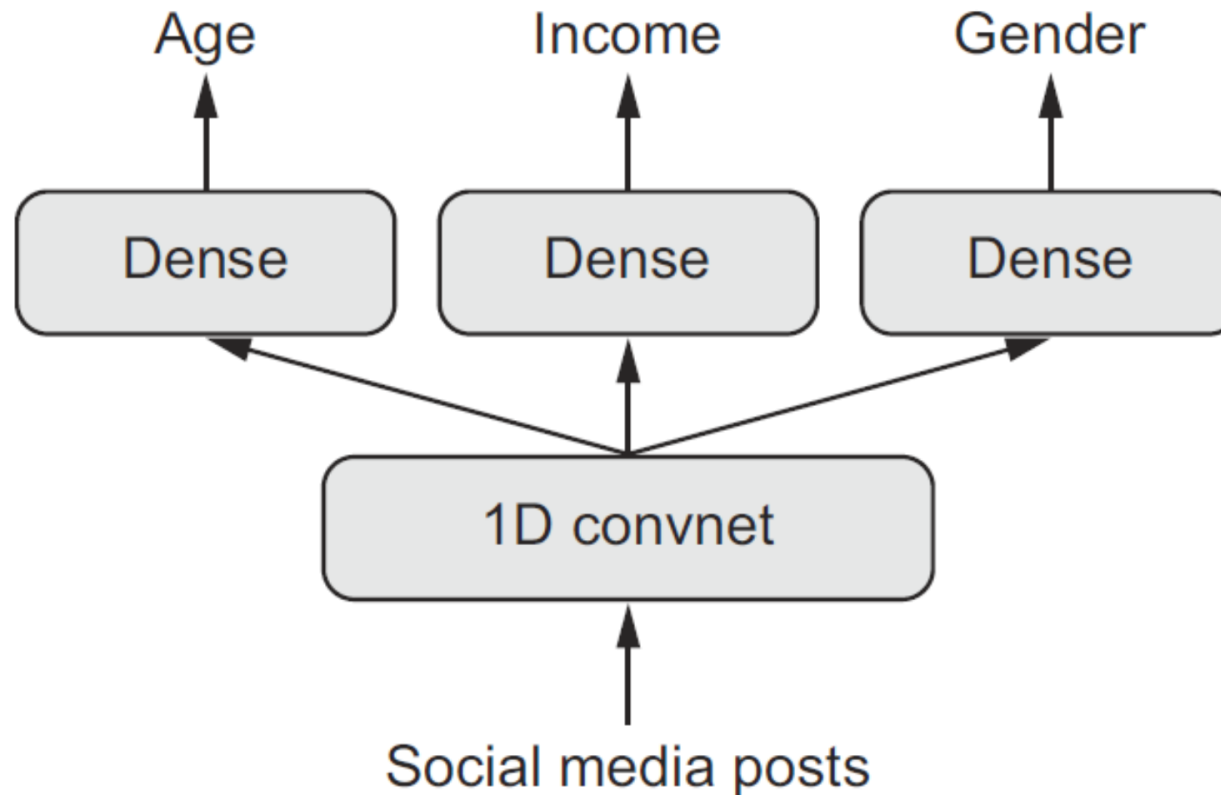
```
model.fit([text, question], answers, epochs=10, batch_size=128)
```

```
### Option 2 ###
```

```
model.fit({'text': text, 'question': question}, answers, epochs=10, batch_size=128)
```

# Multi-output Model

- Predict age, income, gender based on the contents of posts



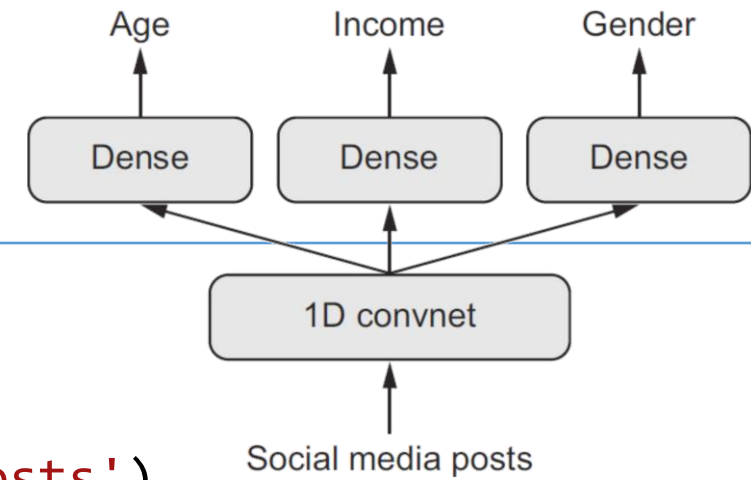
# Multi-output Model

```
vocabulary_size = 50000
num_income_groups = 10
```

```
posts_input = Input(shape=(None,), dtype='int32', name='posts')
embedded_posts = layers.Embedding(256, vocabulary_size)(posts_input)
x = layers.Conv1D(128, 5, activation='relu')(embedded_posts)
x = layers.MaxPooling1D(5)(x)
x = layers.Conv1D(256, 5, activation='relu')(x)
...
x = layers.Dense(128, activation='relu')(x)

age_prediction = layers.Dense(1, name='age')(x)
income_prediction = layers.Dense(num_income_groups, activation='softmax',
                                name='income')(x)
gender_prediction = layers.Dense(1, activation='sigmoid', name='gender')(x)

model = Model(posts_input, [age_prediction, income_prediction, gender_prediction])
```



# Compiling Model with Multiple Losses

```
### Option 1 ###
```

```
model.compile(optimizer='rmsprop',  
              loss=['mse', 'categorical_crossentropy', 'binary_crossentropy'])
```

```
### Option 2 ###
```

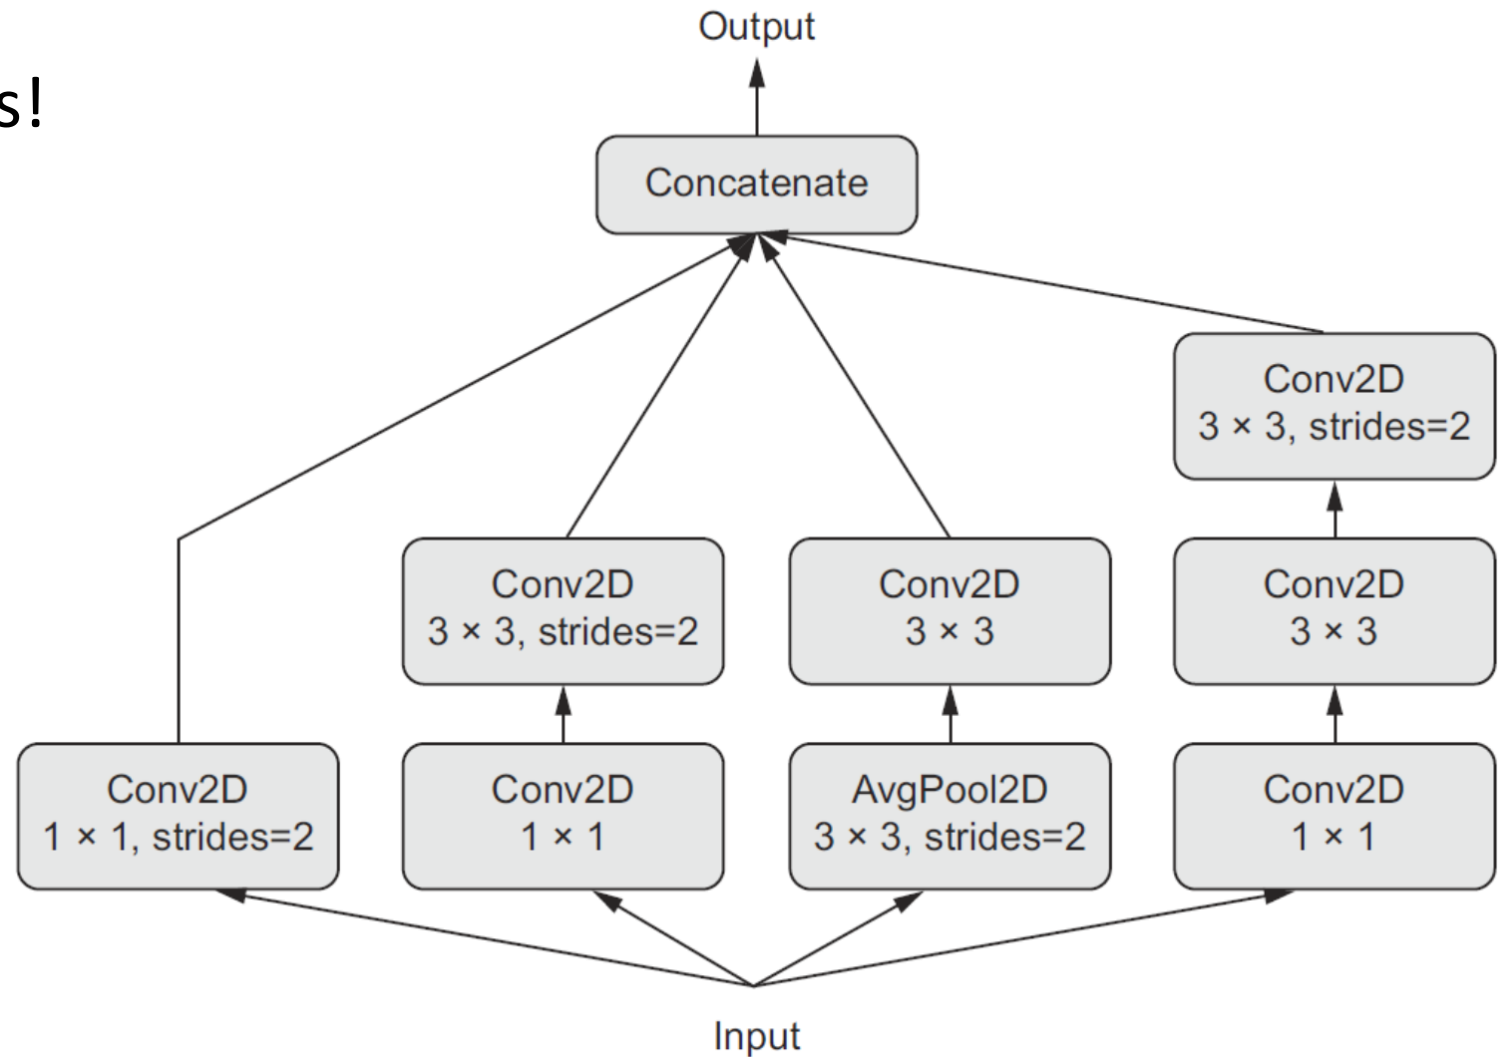
```
model.compile(optimizer='rmsprop',  
              loss={'age': 'mse',  
                    'income': 'categorical_crossentropy',  
                    'gender': 'binary_crossentropy'})
```

# Compile Model with Loss Weighting

```
model.compile(optimizer='rmsprop',  
              loss=['mse', 'categorical_crossentropy', 'binary_crossentropy'],  
              loss_weights=[0.25, 1., 10.])  
  
model.compile(optimizer='rmsprop',  
              loss={'age': 'mse',  
                   'income': 'categorical_crossentropy',  
                   'gender': 'binary_crossentropy'},  
              loss_weights={'age': 0.25,  
                             'income': 1.,  
                             'gender': 10.})
```

# Directed Acyclic Graph of Layers

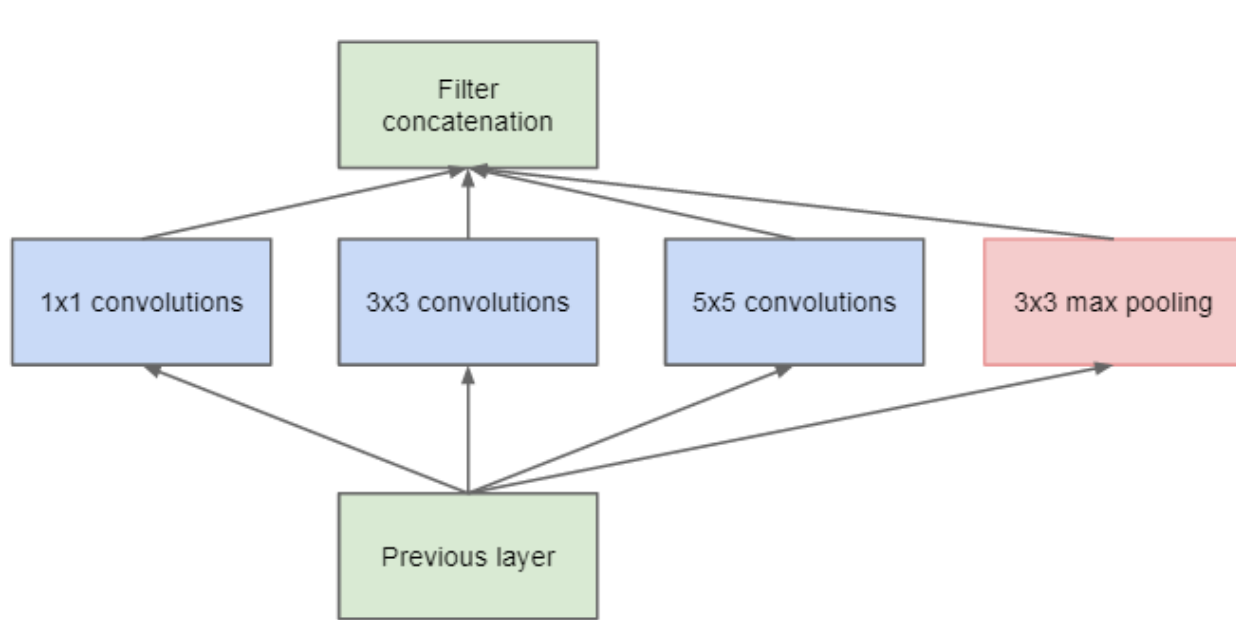
- Graph can't have cycles!



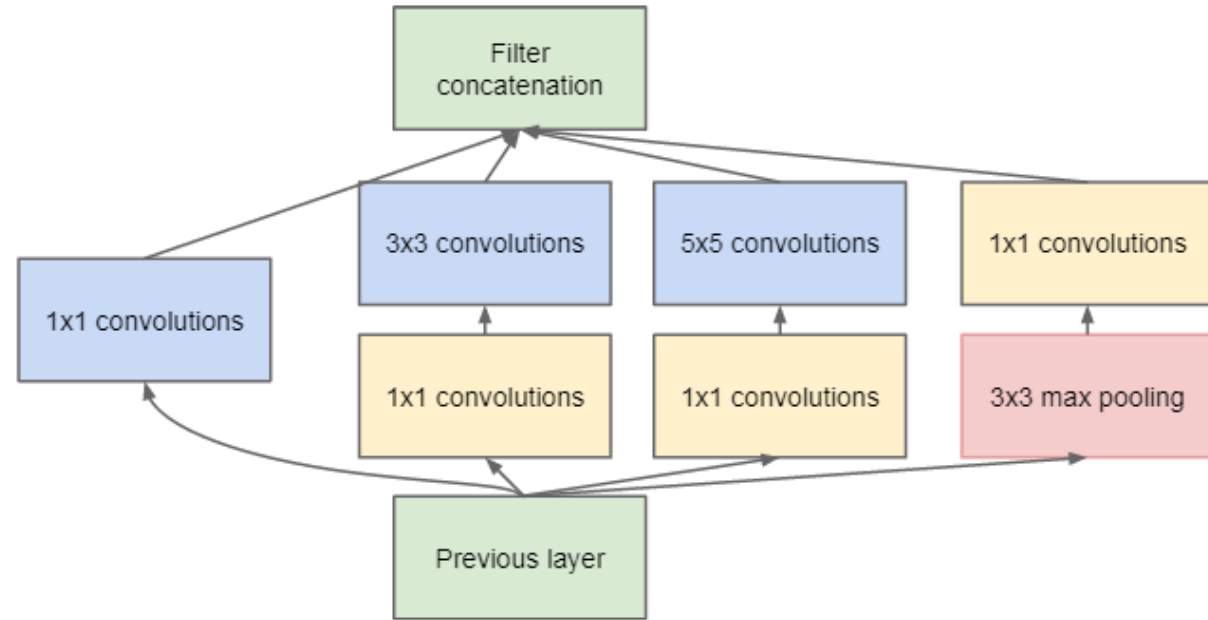


# The Purpose of 1x1 Convolutions

- Reduce the channel dimension



(a) Inception module, naïve version



(b) Inception module with dimension reductions

# Implement Inception Module

```
from keras import layers
branch_a = layers.Conv2D(128, 1, activation='relu', strides=2)(x)

branch_b = layers.Conv2D(128, 1, activation='relu')(x)
branch_b = layers.Conv2D(128, 3, activation='relu', strides=2)(branch_b)

branch_c = layers.AveragePooling2D(3, strides=2)(x)
branch_c = layers.Conv2D(128, 3, activation='relu')(branch_c)

branch_d = layers.Conv2D(128, 1, activation='relu')(x)
branch_d = layers.Conv2D(128, 3, activation='relu')(branch_d)
branch_d = layers.Conv2D(128, 3, activation='relu', strides=2)(branch_d)


output = layers.concatenate([branch_a, branch_b, branch_c, branch_d], axis=-1)
```

# Residual Connection

```
from keras import layers

x = ...
y = layers.Conv2D(128, 3,
                  activation='relu',
                  padding='same')(x)
y = layers.Conv2D(128, 3,
                  activation='relu',
                  padding='same')(y)
y = layers.Conv2D(128, 3,
                  activation='relu',
                  padding='same')(y)

y = layers.add([y, x])
```



# Vanishing Gradients in Deep Learning

- A signal becomes smaller after propagated through multi-layers, and may be lost (vanished)
- Solutions:
  - LSTM: using carry track to propagate signal parallel to main track
  - Residual: simple jump connection

# Share layers and models

- Example: dual-camera

```
from keras import layers
from keras import applications
from keras import Input

xception_base = applications.Xception(weights=None, include_top=False)

left_input = Input(shape=(250, 250, 3))
right_input = Input(shape=(250, 250, 3))
# Extract features from left and right cameras
left_features = xception_base(left_input)
right_input = xception_base(right_input)

merged_features = layers.concatenate([left_features, right_input], axis=-1)
```

# Monitoring Model Training

- Model checkpoint saving
  - Saving the current weights of the model during training
- Early stopping
- Dynamically adjusting parameters
  - Adaptive learning rate during training
- Visualizing the model and data

# Using Callbacks

- **EarlyStopping** - interrupts training when accuracy has stopped improving for more than one epoch
- **ModelCheckpoint** - Saves the current weights after every epoch

```
from keras import callbacks

callbacks_list = [
    callbacks.EarlyStopping(monitor='acc', patience=1),
    callbacks.ModelCheckpoint(filepath='my_model.h5', monitor='val_loss',
                              save_best_only=True)
]

model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.fit(x, y, epochs=10, batch_size=32, callbacks=callbacks_list,
          validation_data=(x_val, y_val))
```



# ReduceLROnPlateau Callback

- factor – the learning rate is multiplied by factor after pre-defined epochs
- patience – epochs before callback is triggered

```
callbacks_list = [  
    callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=10)  
]  
model.fit(x, y, epochs=10, batch_size=32,  
          callbacks=callbacks_list,  
          validation_data=(x_val, y_val))
```

# Implement Your Own Callback Function

- Inherit `keras.callbacks.Callback` and implement any number of the following methods
  - `on_epoch_begin`
  - `on_epoch_end`
  - `on_batch_begin`
  - `on_batch_end`
  - `on_train_begin`
  - `on_train_end`

# Exmaple: Creating Your Own Logger

```
class ActivationLogger(keras.callbacks.Callback):
    def set_model(self, model):
        self.model = model
        layer_outputs = [layer.output for layer in model.layers]
        self.activations_model = keras.models.Model(model.input, layer_outputs)

    def on_epoch_end(self, epoch, logs=None):
        if self.validation_data is None:
            raise RuntimeError('Requires validation_data.')
        validation_sample = self.validation_data[0][0:1]
        activations = self.activations_model.predict(validation_sample)
        f = open('activations_at_epoch_' + str(epoch) + '.npz', 'w')
        np.savez(f, activations)
        f.close()
```

# Tensor Board

- Add `TensorBoard` callback function and assign `log_dir`

```
callbacks = [  
    keras.callbacks.TensorBoard(log_dir='my_log_dir', histogram_freq=1,  
                                embeddings_freq=1)  
]  
  
history = model.fit(x_train, y_train, epochs=20, batch_size=128,  
                    validation_split=0.2, callbacks=callbacks)
```

- Run command => `$ tensorboard --logdir=my_log_dir --host localhost`

# Example: Text Classification with TensorBoard (2-1)

```
import keras
import numpy as np
from keras import layers
from keras.datasets import imdb
from keras.preprocessing import sequence

max_features = 2000
max_len = 500
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)
x_train = sequence.pad_sequences(x_train, maxlen=max_len)
x_test = sequence.pad_sequences(x_test, maxlen=max_len)

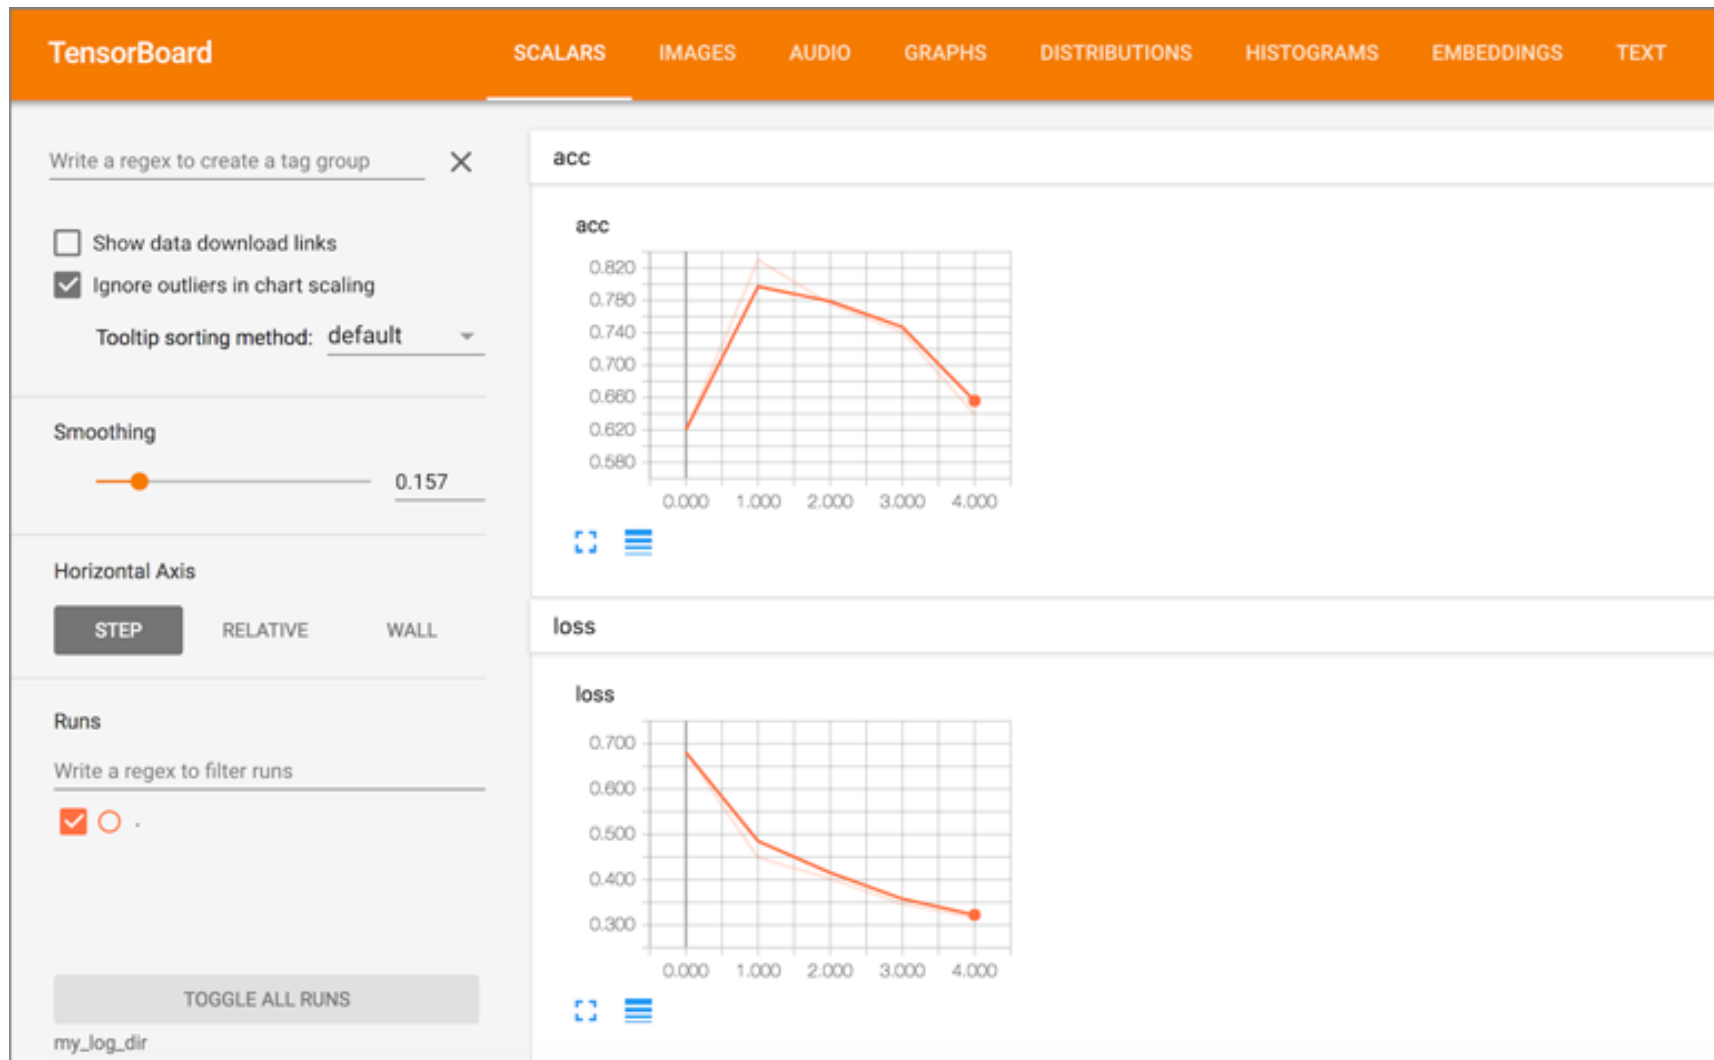
model = keras.models.Sequential()
model.add(layers.Embedding(max_features, 128, input_length=max_len, name='embed'))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.MaxPooling1D(5))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.GlobalMaxPooling1D())
model.add(layers.Dense(1))
model.summary()
```

# Example: Text Classification with TensorBoard (2-2)

```
...
model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['acc'])

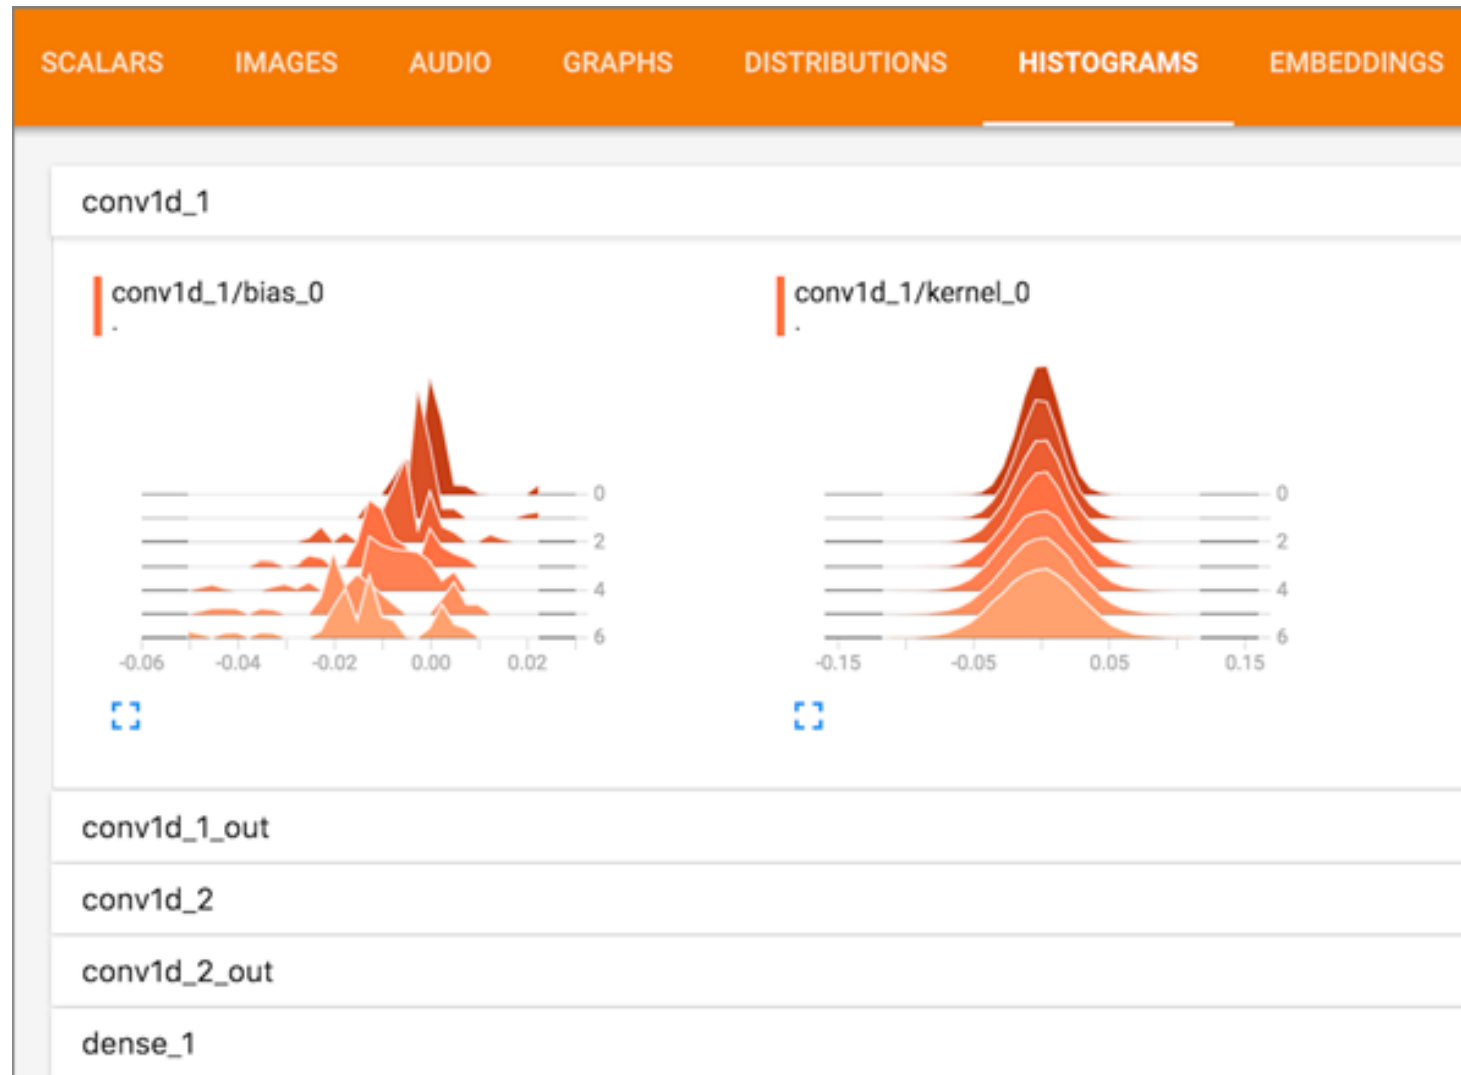
callbacks = [
    keras.callbacks.TensorBoard(
        log_dir='my_log_dir',
        histogram_freq=1,
        embeddings_freq=1,
        embeddings_data = np.arange(0, max_len).reshape((1, max_len)),
    )
]
history = model.fit(x_train, y_train,
                   epochs=20,
                   batch_size=128,
                   validation_split=0.2,
                   callbacks=callbacks)
```

# TensorBoard: Accuracy and Loss

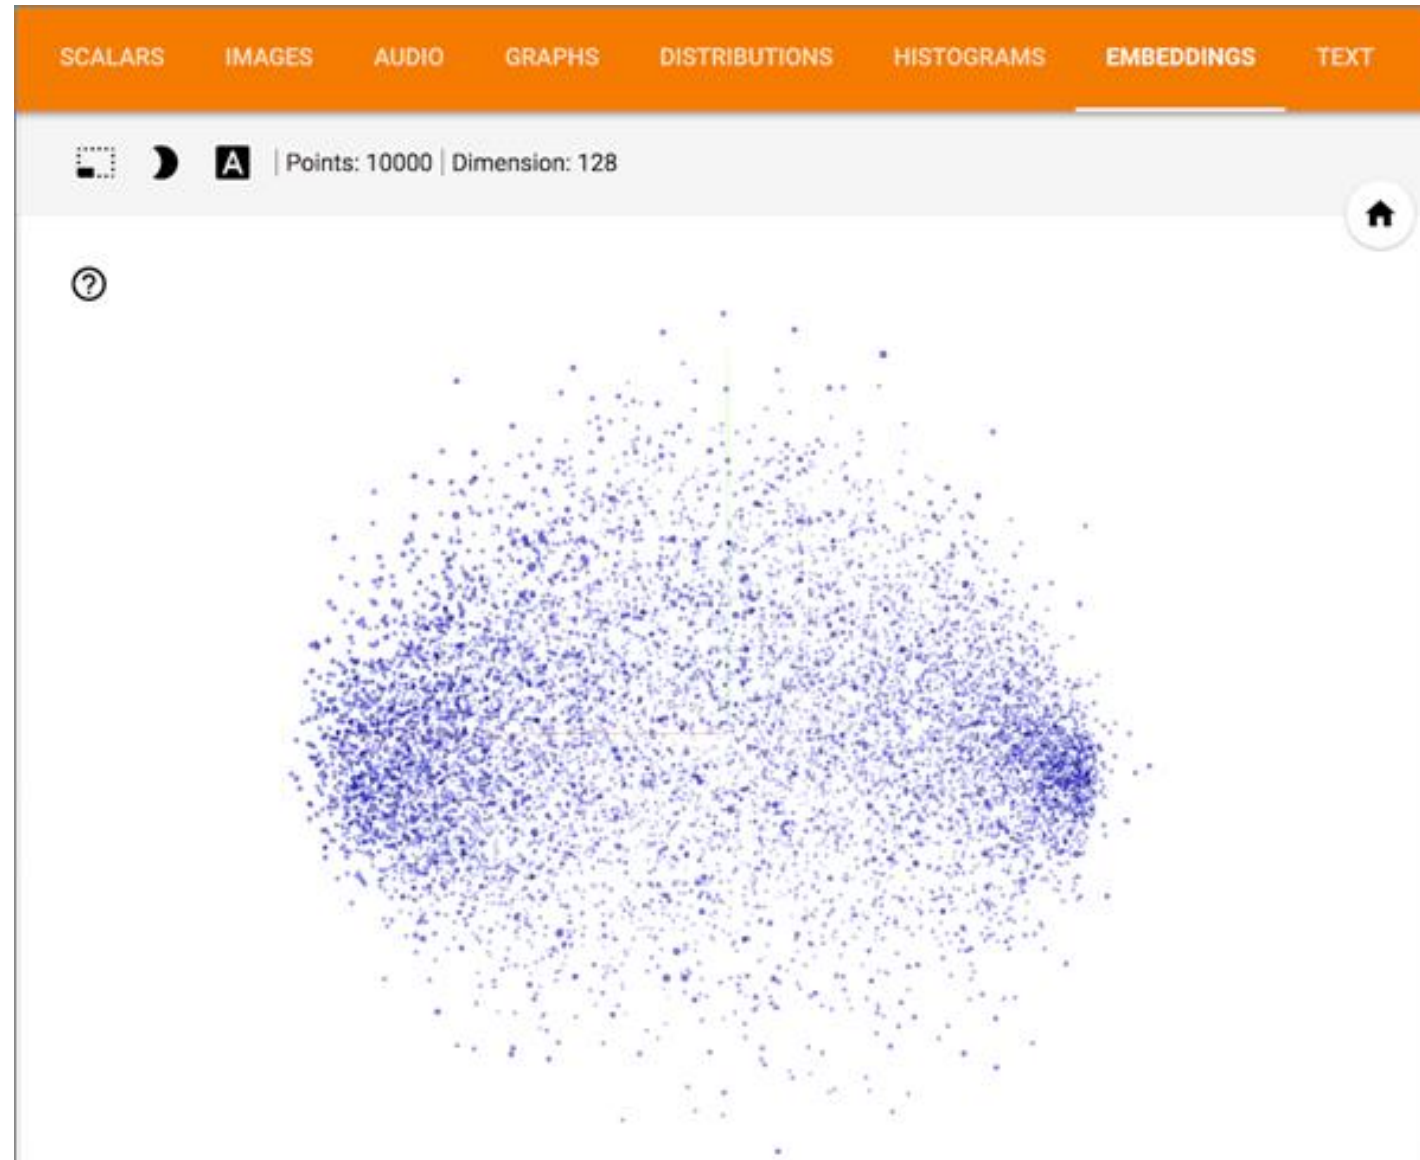




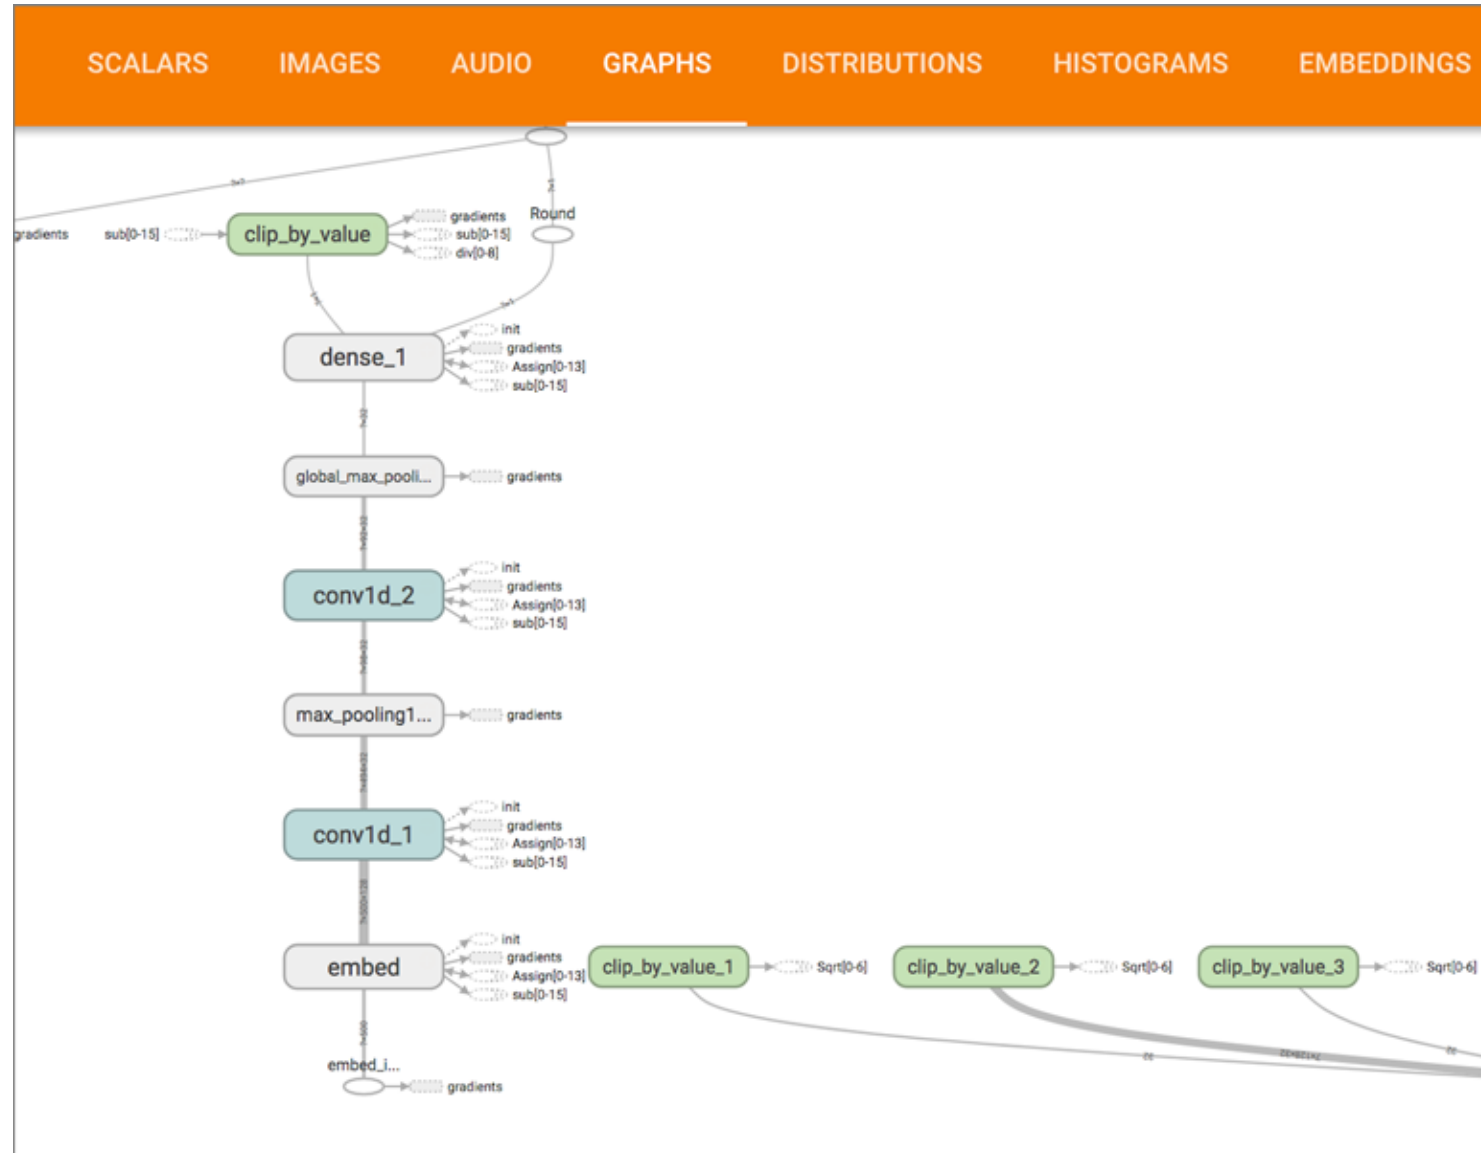
# TensorBoard: Activation Histograms



# TensorBoard: Word-embedding Visualization

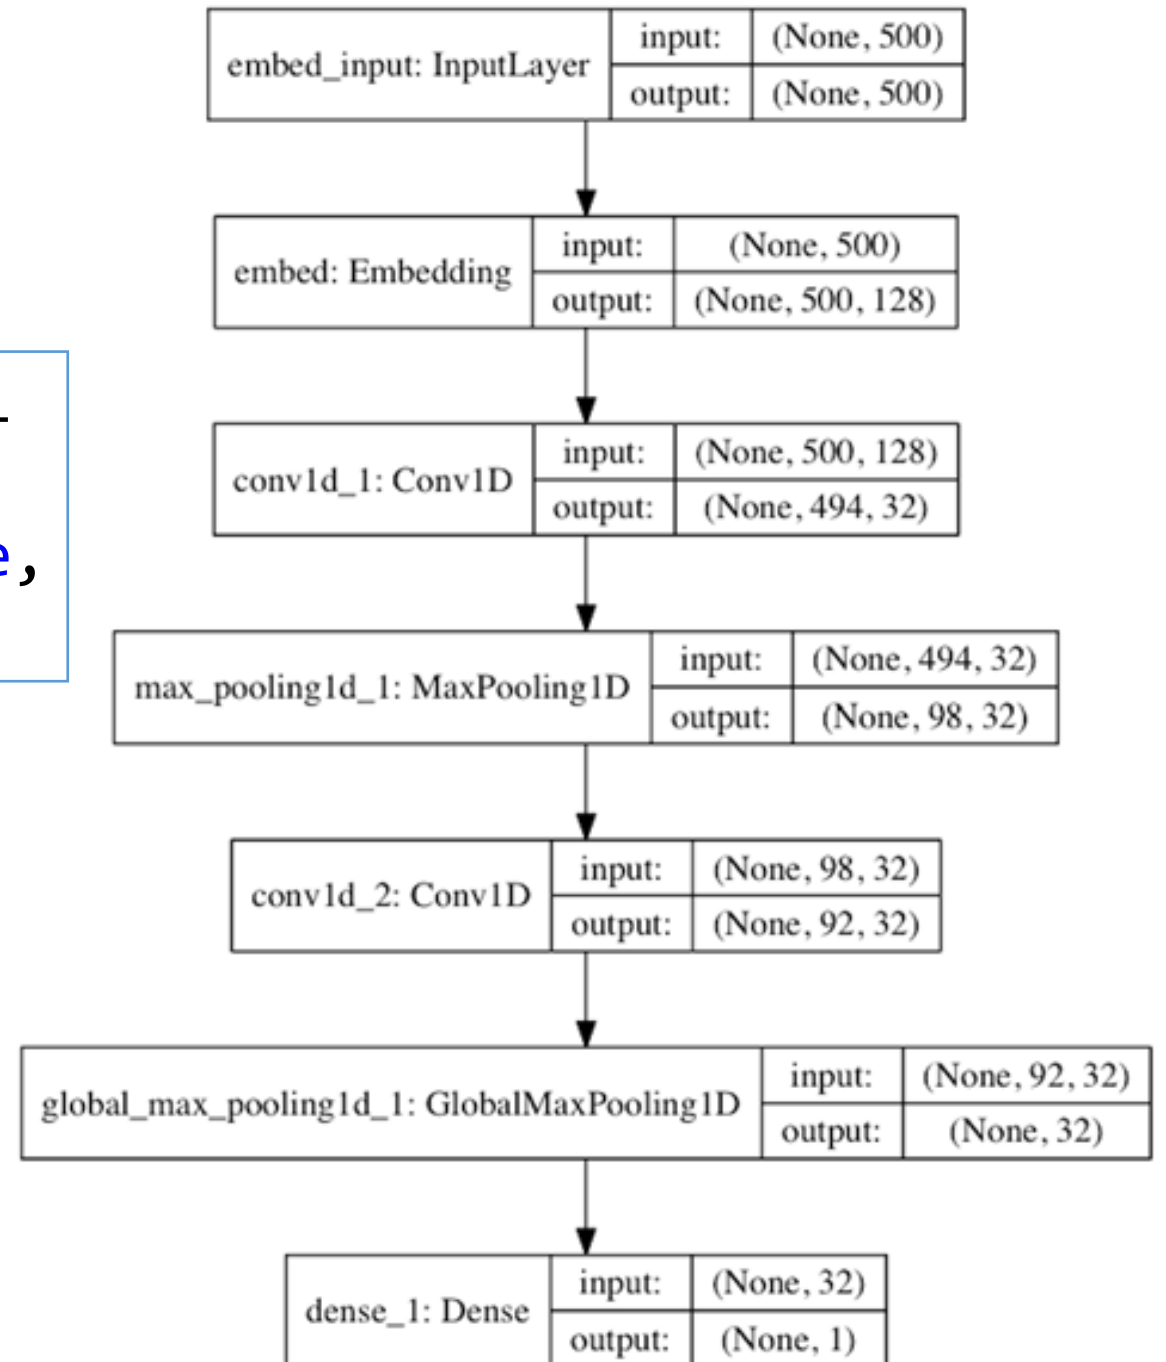


# TensorBoard: Network Graph Visualization



# Keras plot\_model

```
from keras.utils import plot_model  
plot_model(model, show_shapes=True,  
           to_file='model.png')
```



# Batch Normalization

- Sergey Ioffe and Christian Szegedy, “Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift,” *ICML*, 2015 (<https://arxiv.org/abs/1502.03167>).
- Normalizing data after every transformation
- Enhance back propagation
- Some deep networks can only be trained with batch normalization

```
conv_model.add(layers.Conv2D(32, 3, activation='relu'))  
conv_model.add(layers.BatchNormalization())
```

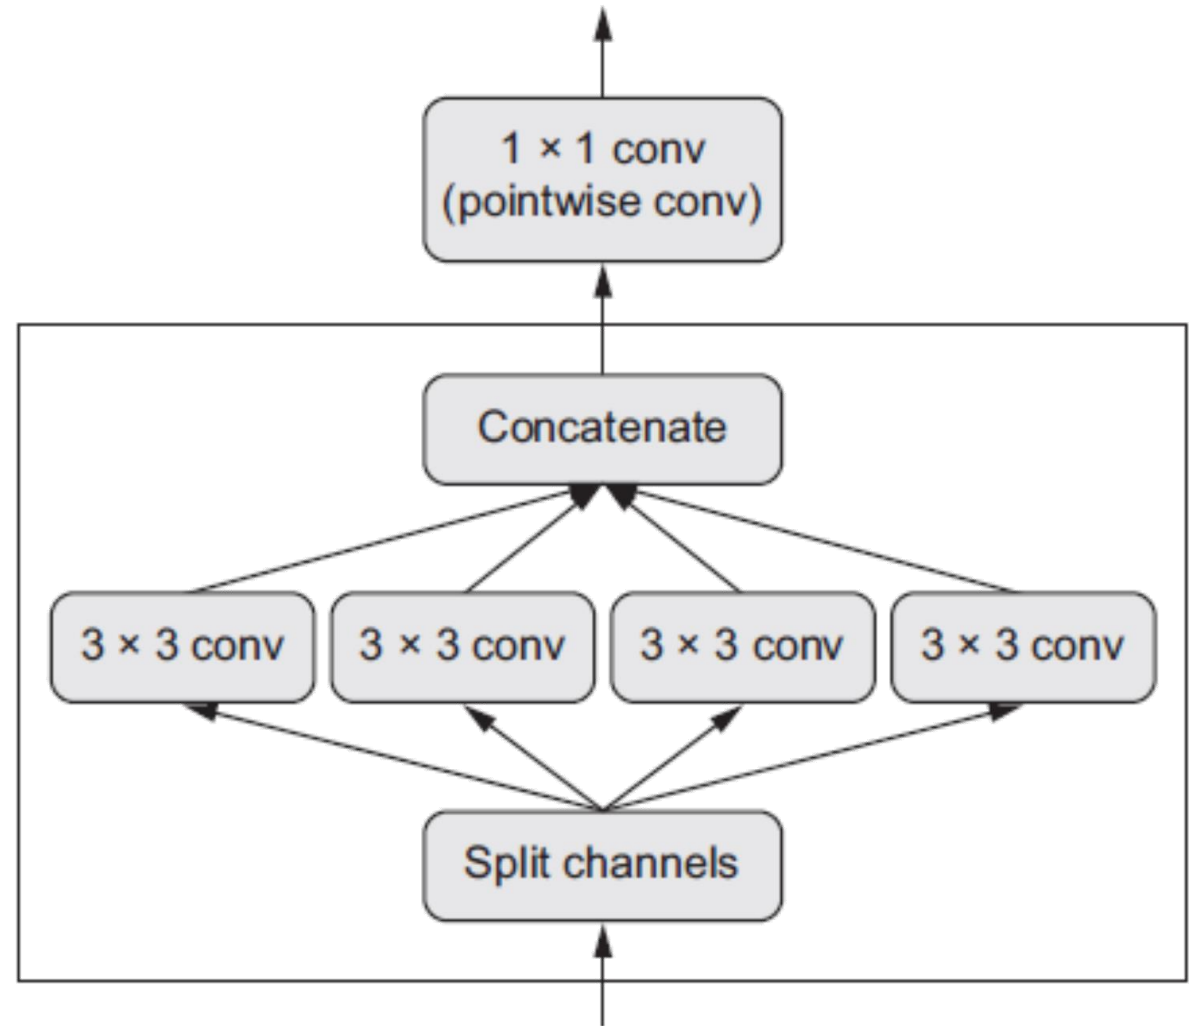
```
dense_model.add(layers.Dense(32, activation='relu'))  
dense_model.add(layers.BatchNormalization())
```

# Batch Renormalization

- Sergey Ioffe, “Batch Renormalization: Towards Reducing Minibatch Dependence in Batch-Normalized Models,” 2017, <https://arxiv.org/abs/1702.03275>.
- Günter Klambauer et al., “Self-Normalizing Neural Networks,” *NIPS*, 2017, <https://arxiv.org/abs/1706.02515> .

# Depthwise Separable Convolution

- Separating the learning of spatial features and channel-wise features
- Less parameters, slightly better accuracy
- Francois Chollet, “Xception: Deep Learning with Depthwise Separable Convolutions.”





# Hyperparameter Optimization

- Random search, genetic algorithm, Bayesian optimization
- Hyperopt (<https://github.com/hyperopt/hyperopt>)
- Hyperas (<https://github.com/maxpumperla/hyperas>)

# Model Ensembling

- Combine the outputs of multiple models (a.k.a late fusion)
  - Random forest
  - Gradient-boosted trees
  - Wide and deep model

# References

- Francois Chollet, “Deep Learning with Python,” Chapter 7